

Prediction!

Contents

Prediction	1
Relating Explanatory Variables in Prediction	4
Fitting a Linear Regression Model in R	8
Another Modeling Approach (k Nearest Neighbors)	18
Competition!	34

Ok, so I think I start off a bit loose to get people comfortable and bring out some playing cards.

- I'll point to someone and ask them to guess the suit of the next card I show.
- Reshuffle and repeat giving them five cards/guesses.
- Then I'll repeat with another three-four students.
- # correctly guessed will be noted somewhere
- They've done a similar thing before but now an emphasis on guessing the next # of correct.

What's the point? How can I best predict the number of card suits the next person will get right?


Prediction

Goal: Predict a new value of a variable

- Ex: Another student will be guessing. Define Y = # of card suits guessed correctly from the five. What should we guess/predict for the next value of Y ?

App

- Chat with them.
- Lead them to talk about a sample.
- Simulate values of Y using an app (in the repo, use `shiny::runGitHub("caryAcademy", username = "jbpost2", subdir = "CardSim", ref = "main").app`)
- Lead them to ideas of using something like the sample mean or median as the predicted value.
- Why something like the sample mean or sample median? What are we really trying to do? Find a value that is 'close' to the most values, i.e. something in the center being the most logical thing to do.



Loss function

Let's assume we have a sample of n people that each guessed five cards. Call these values y_1, y_2, \dots, y_n .

Need: A way to quantify how well our prediction is doing... Suppose there is some best prediction, call it c . How do we measure the quality of c ?

- Using the idea that we want something 'close' to all points, we find a way to compare each point to our prediction. - Think about things like:

$$\begin{aligned} & y_1 - c, (y_1 - c)^2, |y_1 - c| \\ & \sum_{i=1}^n (y_i - c), \sum_{i=1}^n (y_i - c)^2, \sum_{i=1}^n |y_i - c| \\ & \frac{1}{n} \sum_{i=1}^n (y_i - c), \frac{1}{n} \sum_{i=1}^n (y_i - c)^2, \frac{1}{n} \sum_{i=1}^n |y_i - c| \end{aligned}$$

- Quick app to look at how the measures work in the app. **app** - In the end an objective function (mean squared error here) must be created to minimize that uses a 'Loss function', and we'll talk about why we'll use the common squared error loss:

$$g(y_1, \dots, y_n) = \frac{1}{n} \sum_{i=1}^n L(y_i, c) = \frac{1}{n} \sum_{i=1}^n (y_i - c)^2$$

Can we choose an 'optimal' value for c to minimize this function? Calculus to the rescue!

Steps to minimize a function with respect to c :

1. Take the derivative with respect to c
2. Set the derivative equal to 0
3. Solve for c to obtain the potential maximum or minimum
4. Check to see if you have a maximum or minimum (or neither)

Answer comes out to be \bar{y} as the minimizer.

Big wrap: This means that the sample mean is the best prediction when using squared error loss (root mean square error).

Using a Population Distribution

Rather than using sample data, suppose we think about the theoretical distribution for $Y = \#$ of card suits guessed correctly from the five. What might we use here? What assumptions do we need to make this distribution reasonable?

- $Y \sim \text{Bin}(5, 0.25)$ assuming we have independent and identical trials - This gives

$$p_Y(y) = P(Y = y) = \binom{n}{y} p^y (1-p)^{n-y} = \binom{5}{y} 0.25^y 0.75^{5-y}$$

for $y = 0, 1, 2, \dots, n$ or $y = 0, 1, 2, 3, 4, 5$

Is there an optimal value c for the **expected value** of the loss function?

That is, can we minimize (as a function of c) $E[(Y - c)^2]$?

- I think I'll start them out on this one as theory leads to more difficult ideas and I'm not sure if you did general expected values. - In the end though we end up with np or 1.25. - I'll show/discuss that this works generally for any distribution $p_Y(y)$ that has a mean - Discuss the relationship with this and a sample (1/n weight for each point vs $p_Y(y)$ weight for each.) - **Big idea:** This implies that μ is the best predictor to use if you are considering minimizing the expected squared error loss.

HW for after day 1:

- Give them a data set and have them find the mean in R and note that the prediction they would use is that.
- Have them play minesweeper and record the data appropriately. Ask them to produce a best guess for their # of overall bombs.
- Give them a partial derivative question to practice on.

This would be the 2nd (short day) material.

- Recap the big idea of prediction:
 - Need to quantify how well we are doing (squared error loss and MSE)
 - Sample mean is optimal if we have a sample
 - Given a theoretical distribution, expected squared error loss is optimized at the mean of the distribution
- Introduce next material with minesweeper, because it is now browser based, nostalgia on my part, and it seems somewhat fun <https://minesweeper.online/game/938135731>
 - Each student will be assigned a certain number of mines for the board (15x40).
 - They'll click on the first square just below the smiley face. They'll continue to click down one block at a time until they hit a bomb.
 - They'll record in a shared spreadsheet their number of blocks down the first bomb appears.
 - Each person should play 10 games and put their data in.
- Now we'll discuss how we could predict the number of blocks until the first bomb as a function of the number of bombs.
- We'll read in the data to R and do some plotting (I'm not sure what the relationship will be exactly but I'd guess not super linear).

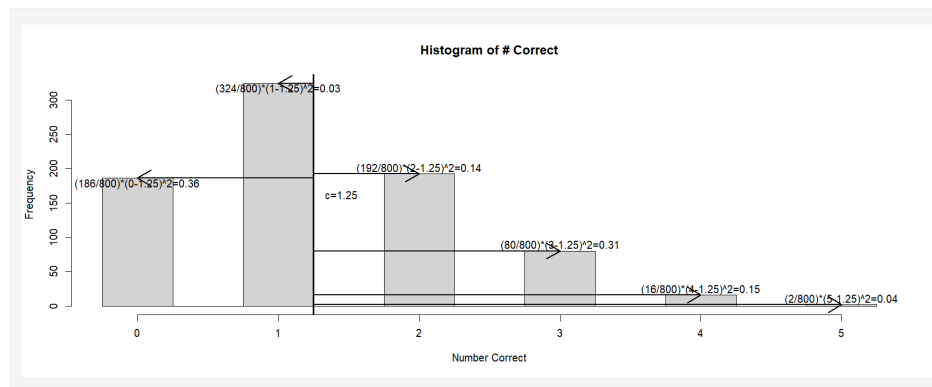
Relating Explanatory Variables in Prediction

Y is a random variable and we'll consider the x values fixed (we'll denote this as $Y|x$). We hope to learn about the relationship between Y and x .

When we considered just Y by itself and used squared error loss, we know that $E(Y) = \mu$ minimizes

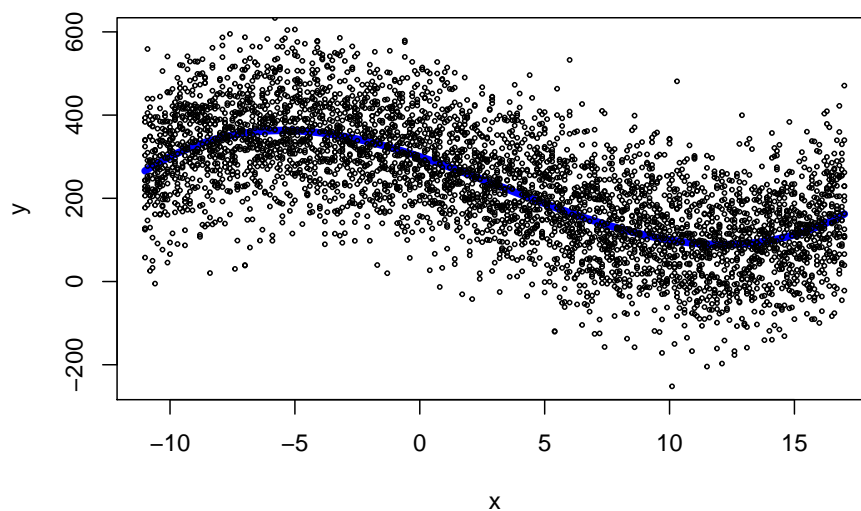
$$E[(Y - c)^2]$$

as a function of c . Given data, we used $\hat{\mu} = \bar{y}$ as our prediction.



Harder (and more interesting) problem is to consider predicting a (response) variable Y as a function of an explanatory variable x .

Below: Blue line, $f(x)$, is the 'true' relationship between x and y



Now that we have an x , $E(Y|x)$ will minimize

$$E[(Y - c)^2|x]$$


We can call this true unknown value $E(Y|x) = f(x)$. That is, the average value of Y will now be considered as a function of x .

Given observed Y 's and x 's, we can estimate this function as $\hat{f}(x)$ (think \bar{y} from before). This $\hat{f}(x)$ will minimize


$$g(y_1, \dots, y_n | x_1, \dots, x_n) = \frac{1}{n} \sum_{i=1}^n L(y_i, \hat{f}(x_i)) = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{f}(x_i))^2$$

Approximating $f(x)$

Although the true relationship is most certainly nonlinear, we may be ok approximating the relationship linearly. For example, consider the same plot as above but between 0 and 5 only:



That's pretty linear. Consider plot between 5 and 15:



Line still does a reasonable job and is often used as a basic approximation.

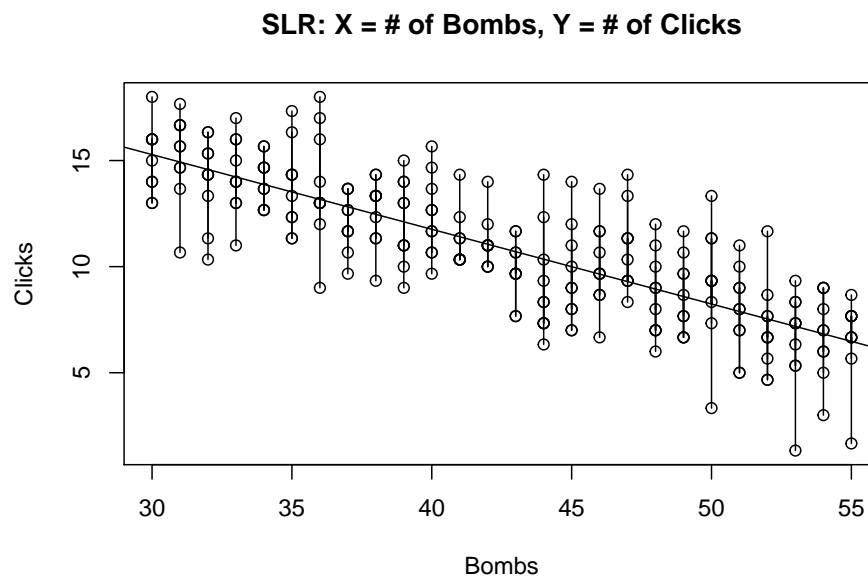
Linear Regression Model

The (fitted) linear regression model uses $\hat{f}(x) = \hat{\beta}_0 + \hat{\beta}_1 x$. This means we want to find the optimal values of $\hat{\beta}_0$ and $\hat{\beta}_1$ from:

$$g(y_1, \dots, y_n | x_1, \dots, x_n) = \sum_{i=1}^n (y_i - \hat{\beta}_0 - \hat{\beta}_1 x_i)^2$$

This equation is often called the 'sum of squared errors (or residuals)' or the 'residual sum of squares'. The model for the data, $E(Y|x) = f(x) = \beta_0 + \beta_1 x$ is called the Simple Linear Regression (SLR) model.

I'll have code at the ready to update and rerender this plot using their data from minesweeper.



Calculus allows us to find the ‘least squares’ estimators, $\hat{\beta}_0$ and $\hat{\beta}_1$ in a nice closed-form!

Do they know partial derivatives? I’m not sure. I think we’ll be running low on time here anyway, so maybe I’ll just talk about the idea of how to get them, set up the equations and then just give the answers.

$$\hat{\beta}_1 = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (x_i - \bar{x})^2}$$

$$\hat{\beta}_0 = \bar{y} - \bar{x}\hat{\beta}_1$$

Then we’ll jump into R and find the values for the minesweeper data and use it to predict by “hand” (plugging it in manually in R).

Ok, day 3 here. Get into R and do some model fitting and predicting. Start with a quick recap here

Fitting a Linear Regression Model in R

Recap: Our goal is to predict a value of Y while including an explanatory variable x . We are assuming we have a sample of (x_i, y_i) pairs, $i = 1, \dots, n$.

The Simple Linear Regression (SLR) model can be used:

$$\hat{f}(x_i) = \hat{y}_i = \hat{\beta}_0 + \hat{\beta}_1 x_i$$

where

- y_i is our response for the i^{th} observation
- x_i is the value of our explanatory variable for the i^{th} observation
- β_0 is the y intercept
- β_1 is the slope

The best model to use if we consider squared error loss has

$$\hat{\beta}_1 = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (x_i - \bar{x})^2}$$
$$\hat{\beta}_0 = \bar{y} - \bar{x}\hat{\beta}_1$$

called the ‘least squares estimates’.

Data Intro

This [dataset](#) contains information about used motorcycles and their cost.

From the information page: This data can be used for a lot of purposes such as price prediction to exemplify the use of linear regression in Machine Learning. The columns in the given dataset are as follows:

- name
- selling price
- year
- seller type
- owner
- km driven
- ex showroom price

The data are available to download from this URL:

<https://www4.stat.ncsu.edu/~online/datasets/bikeDetails.csv>

I think the exploration and modeling would be better to do live and ask them for input as we go through it. I'll put some stuff here and then we can talk about it.

Read in Data and Explore!

```
library(tidyverse)
bikeData <- read_csv("https://www4.stat.ncsu.edu/~online/datasets/bikeDetails.csv")
bikeData <- bikeData %>% tidyr::drop_na()
select(bikeData, selling_price, year, km_driven, ex_showroom_price, name, everything())
```



```
## # A tibble: 626 x 7
##   selling_price year km_driven ex_showroom_price name seller_type owner
##   <dbl> <dbl> <dbl> <dbl> <chr> <chr> <chr>
## 1 150000 2018 12000 148114 Royal Enfi~ Individual 1st ~
## 2 65000 2015 23000 89643 Yamaha Faz~ Individual 1st ~
## 3 18000 2010 60000 53857 Honda CB T~ Individual 1st ~
## 4 78500 2018 17000 87719 Honda CB H~ Individual 1st ~
## 5 50000 2016 42000 60122 Bajaj Disc~ Individual 1st ~
## 6 35000 2015 32000 78712 Yamaha FZ16 Individual 1st ~
## 7 28000 2016 10000 47255 Honda Navi Individual 2nd ~
## 8 80000 2018 21178 95955 Bajaj Aven~ Individual 1st ~
## 9 365000 2019 1127 351680 Yamaha YZF~ Individual 1st ~
## 10 25000 2012 55000 58314 Suzuki Acc~ Individual 1st ~
## # ... with 616 more rows
```

Our 'response' variable here is the `selling_price` and we could use the variable `year`, `km_driven`, or `ex_showroom_price` as the explanatory variable. Let's make some plots and summaries to explore.

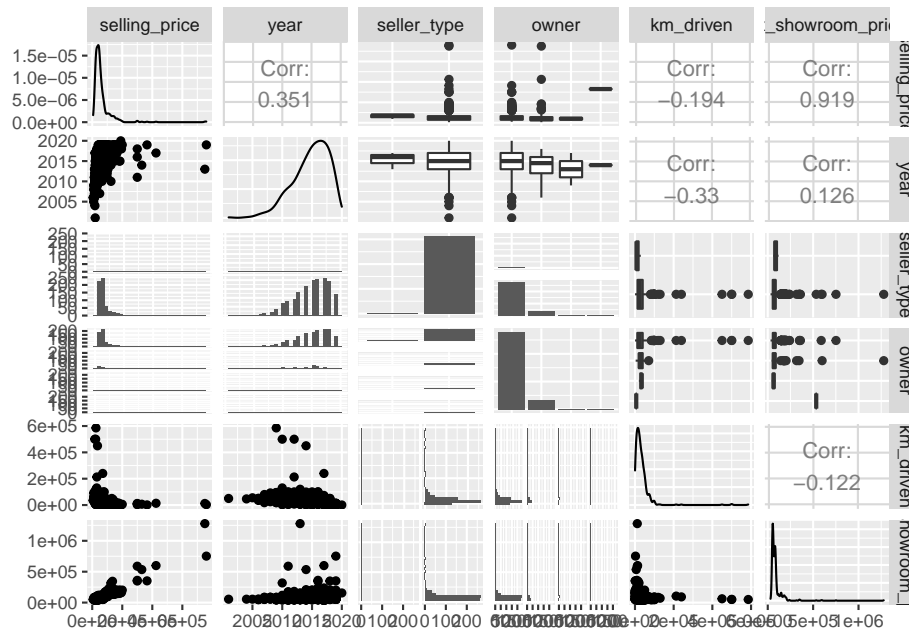
```
summary(bikeData)
```

```
##      name      selling_price      year      seller_type
## Length:626      Min.   : 6000      Min.   :2001      Length:626
## Class :character 1st Qu.: 30000      1st Qu.:2013      Class :character
## Mode  :character Median : 45000      Median :2015      Mode  :character
##                Mean   : 59445      Mean   :2015
##                3rd Qu.: 65000      3rd Qu.:2017
##                Max.   :760000      Max.   :2020
##      owner      km_driven      ex_showroom_price
## Length:626      Min.   : 380      Min.   : 30490
## Class :character 1st Qu.: 13031      1st Qu.: 54852
## Mode  :character Median : 25000      Median : 72753
##                Mean   : 32672      Mean   : 87959
##                3rd Qu.: 40000      3rd Qu.: 87032
##                Max.   :585659      Max.   :1278000
```


```
summarize(group_by(bikeData, owner),
  mean = mean(selling_price),
  median = median(selling_price),
  sd = sd(selling_price),
  IQR = IQR(selling_price))
```

```
## # A tibble: 4 x 5
##   owner      mean median      sd  IQR
##   <chr>    <dbl> <dbl> <dbl> <dbl>
## 1 1st owner 58432. 45000 51125. 35000
## 2 2nd owner 64795. 35000 104861. 30750
## 3 3rd owner 39333. 40000 17010. 17000
## 4 4th owner 330000 330000    NA      0
```

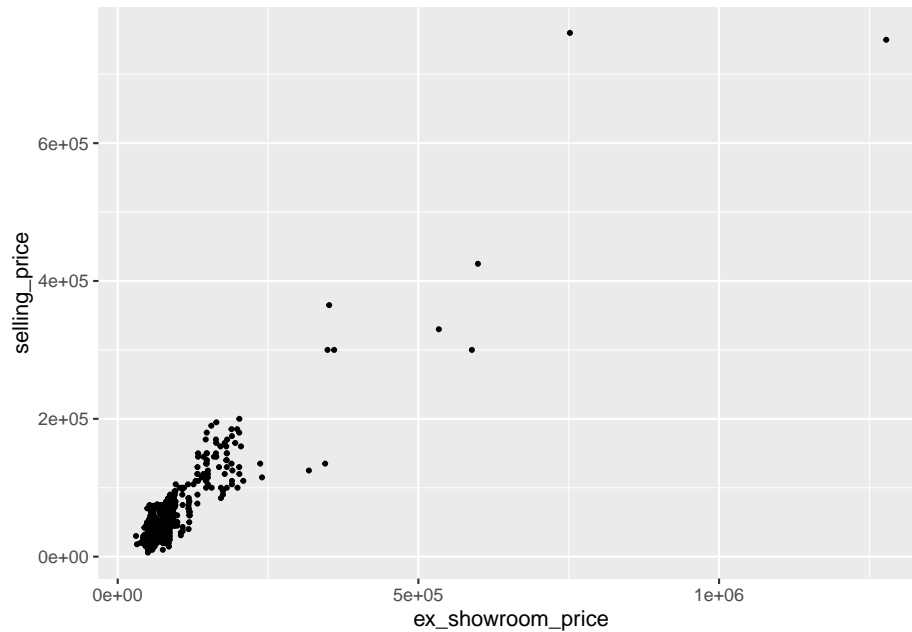
```
library(GGally)
ggpairs(select(bikeData, -name))
```




```
g <- ggplot(data = bikeData, aes(y = selling_price))
g + geom_jitter(aes(x = year, color = km_driven))
```



```
g + geom_point(aes(x = ex_showroom_price), size = 0.75)
```



```
g <- ggplot(data = filter(bikeData, ex_showroom_price < 100000), aes(y = selling_price))
g +
  geom_point(aes(x = ex_showroom_price, color = year, size = km_driven)) +
  scale_color_gradientn(colours = rainbow(5))
```



‘Fitting’ the Model

Basic *linear model* fits done with `lm()`. First argument is a **formula**:

$$\text{response variable} \sim \text{modeling variable}(s)$$

We specify the modeling variable(s) with a `+` sign separating variables. With SLR, we only have one variable on the right hand side.

```
fit <- lm(selling_price ~ ex_showroom_price, data = bikeData)
fit

##
## Call:
## lm(formula = selling_price ~ ex_showroom_price, data = bikeData)
##
## Coefficients:
##      (Intercept)  ex_showroom_price
##      -3010.6984         0.7101
```

We can easily pull off things like the coefficients.

```
coefficients(fit) #helper function

##      (Intercept) ex_showroom_price
##      -3010.6984021         0.7100588
```


Manually predict for an ex_showroom_price of 50000:

```
intercept <- coefficients(fit)[1]
slope <- coefficients(fit)[2]
intercept + slope * 50000
```

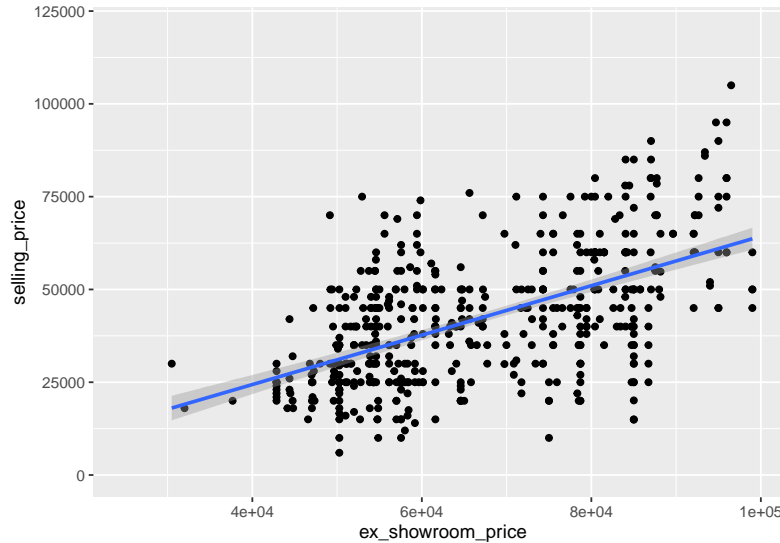
```
## (Intercept)
##      32492.24
```

We can also look at the fit of the line on the graph.

```
ggplot(bikeData, aes(x = ex_showroom_price, y = selling_price)) +
  geom_point() +
  geom_smooth(method = "lm")
```



```
ggplot(bikeData, aes(x = ex_showroom_price, y = selling_price)) +
  geom_point() +
  geom_smooth(method = "lm") +
  scale_x_continuous(limits = c(25000, 100000)) +
  scale_y_continuous(limits = c(0, 120000))
```



Predicting!

Can predict the `selling_price` for a given `ex_showroom_price` easily using the `predict()` function.

```
predict(fit, newdata = data.frame(ex_showroom_price = c(50000, 75000, 100000)))
```

```
##          1          2          3
## 32492.24 50243.71 67995.19
```

Error Assumptions

Although, not needed for prediction, we often assume that we observe our response variable Y as a function of the line plus random errors:

$$Y_i = \beta_0 + \beta_1 x_i + E_i$$

where the errors come from a Normal distribution with mean 0 and variance σ^2 ($E_i \stackrel{iid}{\sim} N(0, \sigma^2)$)

If we do this and use probability theory (maximum likelihood), we will get the same estimates for the slope and intercept as above!

What we get from the normality assumption (if reasonable) is the knowledge of the distribution of our estimators ($\hat{\beta}_0$ and $\hat{\beta}_1$).

What does knowing the distribution allow us to do? We can create confidence intervals or conduct hypothesis tests.

Discuss basic ideas/point of each method.

- Get standard error (SE) for prediction

```
predict(fit, newdata = data.frame(ex_showroom_price = c(50000, 75000, 100000)), se.fit = TRUE)
```

```
## $fit
##          1          2          3
## 32492.24 50243.71 67995.19
##
## $se.fit
##          1          2          3
## 1054.7005  960.2046  958.4166
##
```

```
## $df
## [1] 624
##
## $residual.scale
## [1] 23694.8
```

- Get confidence interval for mean response

```
predict(fit, newdata = data.frame(ex_showroom_price = c(50000, 75000, 100000)),
       se.fit = TRUE, interval = "confidence")
```

```
## $fit
##      fit      lwr      upr
## 1 32492.24 30421.05 34563.44
## 2 50243.71 48358.09 52129.34
## 3 67995.19 66113.07 69877.30
##
## $se.fit
##      1      2      3
## 1054.7005 960.2046 958.4166
##
## $df
## [1] 624
##
## $residual.scale
## [1] 23694.8
```

- Get prediction interval for new response

```
predict(fit, newdata = data.frame(ex_showroom_price = c(50000, 75000, 100000)),
       se.fit = TRUE, interval = "prediction")
```


```
## $fit
##      fit      lwr      upr
## 1 32492.24 -14085.045 79069.53
## 2 50243.71  3674.309 96813.12
## 3 67995.19 21425.922 114564.45
##
## $se.fit
##      1      2      3
## 1054.7005 960.2046 958.4166
##
## $df
## [1] 624
##
## $residual.scale
## [1] 23694.8
```

- Can see the confidence and prediction bands on the plot:

```
library(ciTools)
bikeData <- add_pi(bikeData, fit, names = c("lower", "upper"))

ggplot(bikeData, aes(x = ex_showroom_price, y = selling_price)) +
  geom_point() +
  geom_smooth(method = "lm", fill = "Blue") +
  geom_ribbon(aes(ymin = lower, ymax = upper), alpha = 0.3, fill = "Red") +
```

```
ggtitle("Scatter Plot with 95% PI & 95% CI")
```



For HW have them jump into R and run the code to read in the minesweeper data. Then they could use `lm()` to fit a model and predict.

Multiple Linear Regression

We can add in more than one explanatory variable using the formula for `lm()`. The ideas all follow through!

Just show new SSE type equation.

```
fit <- lm(selling_price ~ ex_showroom_price + year + km_driven, data = bikeData)
fit
```

```
##
## Call:
## lm(formula = selling_price ~ ex_showroom_price + year + km_driven,
##     data = bikeData)
##
## Coefficients:
##      (Intercept)  ex_showroom_price          year      km_driven
##      -9.429e+06      6.863e-01      4.679e+03     -1.053e-02
```

To predict we now need to specify values for all the explanatory variables.

```
data.frame(ex_showroom_price = c(50000, 75000),
            year = c(2010, 2011),
            km_driven = c(15000, 10000))
```

```
##   ex_showroom_price year km_driven
## 1          50000 2010    15000
## 2          75000 2011    10000
```

```
predict(fit, newdata = data.frame(ex_showroom_price = c(50000, 75000),
                                   year = c(2010, 2011),
```

```
km_driven = c(15000, 10000)),
se.fit = TRUE, interval = "confidence")
```

```
## $fit
##      fit      lwr      upr
## 1 11118.83 7914.815 14322.85
## 2 33007.56 30202.482 35812.63
##
## $se.fit
##      1      2
## 1631.552 1428.402
##
## $df
## [1] 622
##
## $residual.scale
## [1] 19011.31
```

Difficult to visualize the model fit though!

Evaluating Model Accuracy

Which model is better? Ideally we want a model that can predict **new** data better, not the data we've already seen. We need a **test** set to predict on. We also need to quantify what we mean by better!

Training and Test Sets

We can split the data into a **training** set and **test** set.



- On the training set we can fit (or train) our models. The data from the test set isn't used at all in this process.
- We can then predict for the test set observations (for the combinations of explanatory variables seen in the test set). Can then compare the predicted values to the actual observed responses from the test set.

Let's jump into R and fit our SLR model and compare it to an MLR model.

Split data randomly:

```
set.seed(1)
numObs <- nrow(bikeData)
index <- sample(1:numObs, size = 0.7*numObs, replace = FALSE)
train <- bikeData[index, ]
test <- bikeData[-index, ]
```

Fit the models on the training data only.


```
fitSLR <- lm(selling_price ~ ex_showroom_price , data = train)
fitMLR <- lm(selling_price ~ ex_showroom_price + year + km_driven, data = train)
```

Predict on the test set.

```
predSLR <- predict(fitSLR, newdata = test)
predMLR <- predict(fitMLR, newdata = test)
tibble(predSLR, predMLR, test$selling_price)
```

```
## # A tibble: 188 x 3
##   predSLR predMLR `test$selling_price`
##   <dbl>    <dbl>          <dbl>
## 1 100749. 113099.          150000
## 2  59739.  60915.           65000
## 3  58390.  72928.           78500
## 4  39035.  45467.           50000
## 5  52073.  53538.           35000
## 6  64166.  78343.           80000
## 7  79576.  57387.           40000
## 8 100749. 112955.          150000
## 9  89924. 102497.          120000
## 10 36247.  25302.           25000
## # ... with 178 more rows
```

Root Mean Square Error

Which is better?? Can use squared error loss to evaluate! (Square root of the mean squared error loss is often reported instead and is called RMSE or Root Mean Square Error.)

```
sqrt(mean((predSLR - test$selling_price)^2))
```

```
## [1] 27840.6
```

```
sqrt(mean((predMLR - test$selling_price)^2))
```

```
## [1] 23374.61
```

MLR fit does much better at predicting!

Another Modeling Approach (k Nearest Neighbors)

Day 4:

kNN relate to idea with minesweeper data about how they predicted for their individual piece.

Will need to use $k = 10, 30, 50, \dots$

Intro and Recap

Recap: Our previous goal was to predict a value of Y while including an explanatory variable x . With that x , we said $E(Y|x)$ will minimize

$$E[(Y - c)^2|x]$$

We called this true unknown value $E(Y|x) = f(x)$.

Given observed Y 's and x 's, we can estimate this function as $\hat{f}(x)$ (with SLR we estimated it with $\hat{\beta}_0 + \hat{\beta}_1 x$). This $\hat{f}(x)$ will minimize

$$g(y_1, \dots, y_n | x_1, \dots, x_n) = \frac{1}{n} \sum_{i=1}^n L(y_i, \hat{f}(x_i)) = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{f}(x_i))^2$$

What other things could we consider for $f(x)$???

Consider the minesweeper data we collected previously.

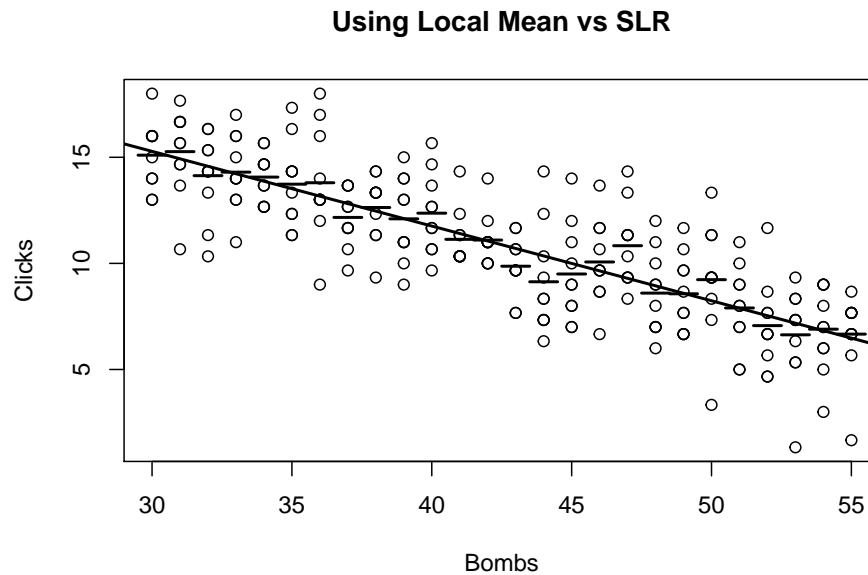
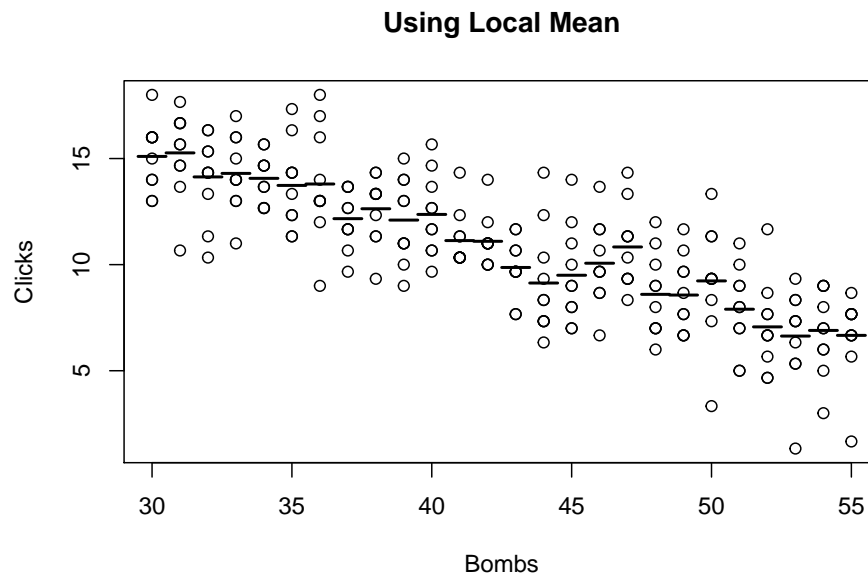
Recall the minesweeper data we collected. In your homework, you were asked to determine a prediction of the number of blocks you could click before hitting a bomb for **your** given number of bombs. You were each creating your estimate of $f(x)$ for your x value!

Let's visualize that idea and compare it to the SLR fit!

I'd need the data here but this is what it would look like. I won't show the code here.

```
## `summarise()` ungrouping output (override with `.groups` argument)

## # A tibble: 26 x 2
##   bombs mean
##   <int> <dbl>
## 1     30 15.1
## 2     31 15.3
## 3     32 14.1
## 4     33 14.3
## 5     34 14.1
## 6     35 13.7
## 7     36 13.8
## 8     37 12.2
## 9     38 12.6
## 10    39 12.1
## # ... with 16 more rows
```



This is the idea of k Nearest Neighbors (kNN) for predicting a numeric response!

kNN

To predict a value of our (numeric) response kNN uses the **average of the k ‘closest’ responses**. For numeric data, we usually use Euclidean distance ($d(x_1, x_2) = \sqrt{(x_1 - x_2)^2}$) to determine the closest values.

- Large k implies more rigid (possibly *underfit* but lower variance prediction).
- Smaller k implies less rigid (possible *overfit* with high variance in prediction)

[Let's check out this app.](#)

For the minesweeper data, we had many values at the same x (# of bombs). That's why we considered using

only 10, 30, 50, ... Otherwise, we have ties and then things get tricky!

Choosing the Value of k

How do we choose which k value to use? We can do a similar training vs test set idea. Fit the models (one model for each k) and predict on the test set. The model with the lowest Root Mean Squared Error (RMSE) on the test set can be chosen!

kNN Models for selling_price from the Bike Dataset

Previously, we fit the SLR model using the `ex_showroom_price` to predict our `selling_price` of motorcycles. We'll refit this using the training data here.

```
fitSLR <- lm(selling_price ~ ex_showroom_price, data = train)
fitSLR


##
## Call:
## lm(formula = selling_price ~ ex_showroom_price, data = train)
##
## Coefficients:
##      (Intercept)  ex_showroom_price
##      -3132.4598           0.7014

summary(fitSLR)

##
## Call:
## lm(formula = selling_price ~ ex_showroom_price, data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -109968  -11974   -1779   10033  121479
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -3.132e+03  1.804e+03  -1.736   0.0832 .
## ex_showroom_price  7.014e-01  1.717e-02  40.849   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 21700 on 436 degrees of freedom
## Multiple R-squared:  0.7928, Adjusted R-squared:  0.7924
## F-statistic: 1669 on 1 and 436 DF, p-value: < 2.2e-16

ggplot(train, aes(x = ex_showroom_price, y = selling_price)) +
  geom_point() +
  geom_smooth(method = "lm")

## `geom_smooth()` using formula 'y ~ x'
```



Obtain the prediction on the test set.

```
predSLR <- predict(fitSLR, newdata = test)
```

Let's now fit the kNN model using a few values of k .

$k = 1$:

```
library("caret")
```

```
## Warning: package 'caret' was built under R version 3.5.3
```

```
## Loading required package: lattice
```

```
##
```

```
## Attaching package: 'caret'
```

```
## The following object is masked from 'package:purrr':
```

```
##
```

```
## lift
```

```
k <- 1
```

```
kNNFit1 <- train(selling_price ~ ex_showroom_price,
  data = train,
  method = "knn",
  tuneGrid = data.frame(k = k)
)
```

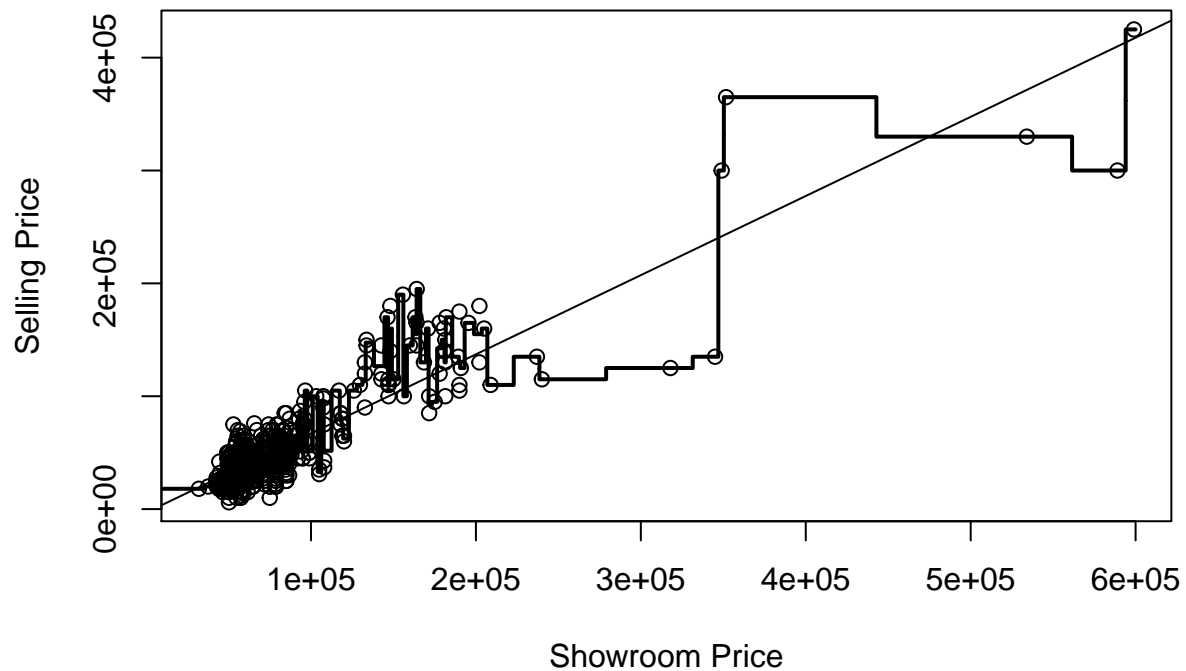
```
kNNFit1
```

```
## k-Nearest Neighbors
```


```
##
```

```
## 438 samples
## 1 predictor
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 438, 438, 438, 438, 438, 438, ...
## Resampling results:
##
## RMSE      Rsquared  MAE
## 23995.06  0.7606492  15197.3
##
## Tuning parameter 'k' was held constant at a value of 1
predkNN1 <- predict(kNNFit1, newdata = test)
```

kNN predictions vs SLR



kNN predictions vs SLR




$k = 10$:

```
k <- 10
kNNFit10 <- train(selling_price ~ ex_showroom_price,
  data = train,
  method = "knn",
  tuneGrid = data.frame(k = k)
)
kNNFit10

## k-Nearest Neighbors
##
## 438 samples
## 1 predictor
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 438, 438, 438, 438, 438, 438, ...
## Resampling results:
##
## RMSE      Rsquared    MAE
## 27190.57  0.7274929  16045.91
##
## Tuning parameter 'k' was held constant at a value of 10
```

```
predkNN10 <- predict(kNNFit10, newdata = test)
```

kNN predictions vs SLR



$k = 20$:


```
k <- 20
kNNFit20 <- train(selling_price ~ ex_showroom_price,
  data = train,
  method = "knn",
  tuneGrid = data.frame(k = k)
)
kNNFit20
```

```
## k-Nearest Neighbors
##
## 438 samples
## 1 predictor
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 438, 438, 438, 438, 438, 438, ...
## Resampling results:
##
## RMSE      Rsquared    MAE
## 24376.62  0.7327945  14873.03
##
## Tuning parameter 'k' was held constant at a value of 20
```



```
predkNN20 <- predict(kNNFit20, newdata = test)
```

kNN predictions vs SLR



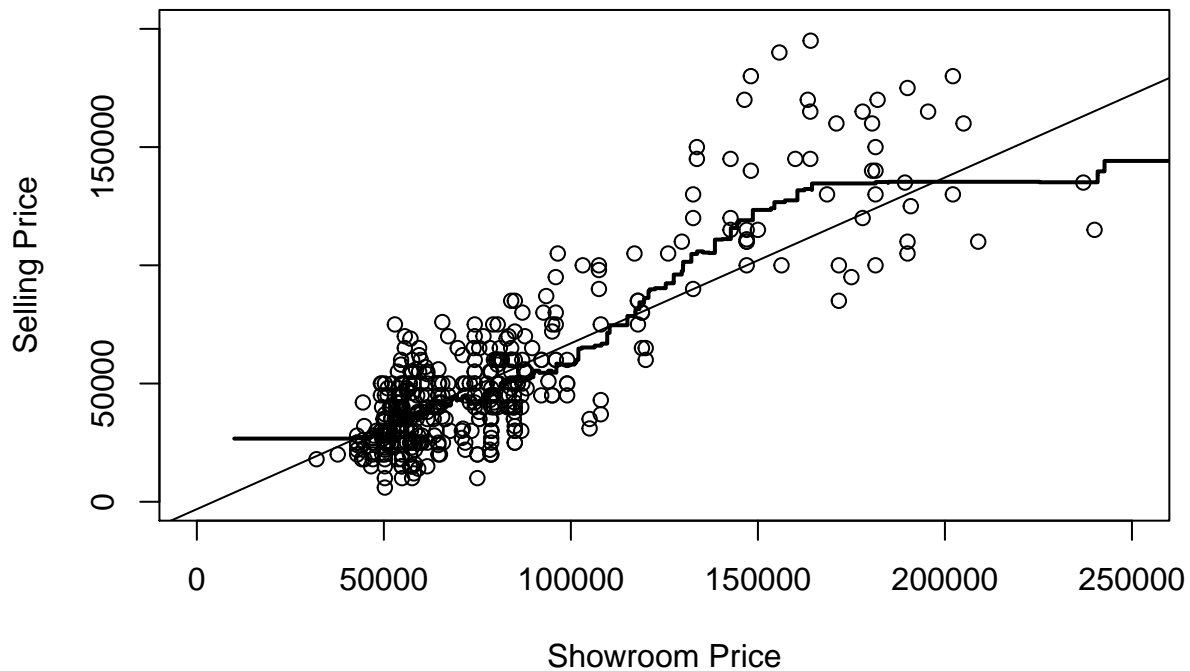
$k = 50$:

```
k <- 50
kNNFit50 <- train(selling_price ~ ex_showroom_price,
  data = train,
  method = "knn",
  tuneGrid = data.frame(k = k)
)
kNNFit50

## k-Nearest Neighbors
##
## 438 samples
## 1 predictor
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 438, 438, 438, 438, 438, 438, ...
## Resampling results:
##
##   RMSE      Rsquared   MAE
## 25970.85  0.7269515 15167.21
##
## Tuning parameter 'k' was held constant at a value of 50
```

```
predkNN50 <- predict(kNNFit50, newdata = test)
```

kNN predictions vs SLR




$k = 100$:

```
k <- 100
kNNFit100 <- train(selling_price ~ ex_showroom_price,
  data = train,
  method = "knn",
  tuneGrid = data.frame(k = k)
)
kNNFit100

## k-Nearest Neighbors
##
## 438 samples
## 1 predictor
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 438, 438, 438, 438, 438, 438, ...
## Resampling results:
##
## RMSE      Rsquared    MAE
## 32864.01  0.7052464  17514.27
##
## Tuning parameter 'k' was held constant at a value of 100
```

```
predkNN100 <- predict(kNNFit100, newdata = test)
```

kNN predictions vs SLR



Compare test set RMSE!

```
RMSE <- function(pred, test){sqrt(mean((pred-test)^2))}
SLR <- RMSE(predSLR, test$selling_price)
kNN1 <- RMSE(predkNN1, test$selling_price)
kNN10 <- RMSE(predkNN10, test$selling_price)
kNN20 <- RMSE(predkNN20, test$selling_price)
kNN50 <- RMSE(predkNN50, test$selling_price)
kNN100 <- RMSE(predkNN100, test$selling_price)

data.frame(method = c("SLR", "kNN1", "kNN10", "kNN20", "kNN50", "kNN100"),
           RMSE = c(SLR, kNN1, kNN10, kNN20, kNN50, kNN100))
```

```
##  method    RMSE
## 1    SLR 27840.60
## 2   kNN1 38426.66
## 3  kNN10 57192.25
## 4  kNN20 61807.92
## 5  kNN50 65410.10
## 6 kNN100 70511.86
```

Ok, of course we don't want to do this manually in real life... What we actually do:

1. Split the data into training and test sets 2. Choose a 'best' model for a given method (MLR, kNN, etc.) using the training set - This requires us to have a method to choose using only the training data! 3. Compare the best model from each method on the test set to see how they do 4. Refit the chosen model on the full

data set. This model would then be what you would use for future predictions.

R makes it easy! To choose a kNN model we can run code like this:

```
k <- 1:100
kNNFit <- train(selling_price ~ ex_showroom_price,
  data = train,
  method = "knn",
  tuneGrid = data.frame(k = k),
  trControl = trainControl(method = "cv", number = 10)
)
kNNFit
```

```
## k-Nearest Neighbors
##
## 438 samples
## 1 predictor
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 393, 395, 395, 394, 395, 394, ...
## Resampling results across tuning parameters:
##
##  k      RMSE      Rsquared    MAE
##  1  21791.19  0.8255668  14210.50
##  2  20782.31  0.8191957  13937.01
##  3  20995.61  0.8143527  13817.78
##  4  20749.55  0.8173147  13758.75
##  5  21038.07  0.8142272  14035.86
##  6  21510.43  0.8097114  14467.62
##  7  22098.16  0.8042000  14627.56
##  8  22504.33  0.8005286  14669.61
##  9  22877.16  0.7966801  14785.62
## 10  23200.59  0.7898105  14802.47
## 11  23582.20  0.7827674  14941.16
## 12  23537.03  0.7831586  14919.37
## 13  23592.47  0.7837519  14842.92
## 14  23736.89  0.7818278  14865.71
## 15  24103.96  0.7746132  14924.28
## 16  24088.63  0.7759797  14859.08
## 17  24185.28  0.7746521  14922.38
## 18  24247.40  0.7735124  14975.27
## 19  24354.65  0.7717801  15033.22
## 20  24590.25  0.7676895  15121.85
## 21  24593.15  0.7672238  15079.50
## 22  24631.29  0.7653932  15041.76
## 23  24735.65  0.7629428  15007.55
## 24  24768.37  0.7628850  14961.07
## 25  24728.38  0.7643760  14913.60
## 26  24720.40  0.7647586  14963.02
## 27  24648.75  0.7682739  14921.69
## 28  24733.92  0.7668029  14992.23
## 29  24768.96  0.7667417  14926.99
## 30  24790.60  0.7668187  15002.75
## 31  24794.68  0.7682305  14984.37
```

##	32	24813.57	0.7698235	14965.24
##	33	24903.82	0.7696879	15054.05
##	34	24881.14	0.7704787	15121.08
##	35	24936.45	0.7691981	15131.56
##	36	24970.41	0.7686061	15129.71
##	37	24995.03	0.7676471	15134.50
##	38	25022.80	0.7677180	15114.04
##	39	25175.24	0.7651497	15151.41
##	40	25213.43	0.7645103	15189.17
##	41	25262.68	0.7638752	15217.27
##	42	25336.92	0.7636167	15266.03
##	43	25398.05	0.7630132	15309.07
##	44	25386.27	0.7637307	15270.05
##	45	25431.92	0.7642146	15291.61
##	46	25520.34	0.7639196	15352.53
##	47	25530.97	0.7666179	15284.97
##	48	25603.50	0.7667640	15266.34
##	49	25617.05	0.7691948	15294.00
##	50	25762.90	0.7674965	15369.59
##	51	25825.95	0.7691776	15394.02
##	52	25934.43	0.7688300	15396.09
##	53	26058.52	0.7680946	15446.35
##	54	26246.79	0.7653493	15457.64
##	55	26372.74	0.7658096	15510.02
##	56	26444.81	0.7653888	15533.50
##	57	26659.06	0.7635288	15646.53
##	58	26626.82	0.7665449	15610.01
##	59	26793.17	0.7644656	15664.50
##	60	26918.21	0.7655529	15729.03
##	61	27039.12	0.7645784	15750.49
##	62	27101.31	0.7659216	15779.22
##	63	27327.46	0.7655766	15898.79
##	64	27364.95	0.7683210	15892.68
##	65	27405.19	0.7704313	15890.13
##	66	27581.50	0.7681817	15982.45
##	67	27769.85	0.7673199	16093.19
##	68	27908.39	0.7671792	16148.90
##	69	28048.51	0.7637428	16188.47
##	70	28142.57	0.7636680	16232.98
##	71	28283.13	0.7619680	16294.08
##	72	28381.09	0.7634655	16324.80
##	73	28457.10	0.7649634	16385.45
##	74	28564.43	0.7647628	16447.17
##	75	28728.52	0.7612627	16506.13
##	76	28829.53	0.7619253	16579.31
##	77	28933.02	0.7621430	16672.18
##	78	29004.64	0.7622336	16710.35
##	79	29180.24	0.7611944	16772.38
##	80	29303.24	0.7610363	16824.51
##	81	29412.15	0.7608901	16889.26
##	82	29494.77	0.7614794	16927.23
##	83	29605.61	0.7612174	16966.20
##	84	29691.34	0.7616438	17030.23
##	85	29825.79	0.7595197	17084.52

```
##      86 29967.33 0.7591394 17160.67
##      87 30082.33 0.7596716 17220.50
##      88 30237.84 0.7579103 17312.90
##      89 30345.25 0.7580539 17374.06
##      90 30622.68 0.7584970 17519.13
##      91 30762.00 0.7586696 17609.22
##      92 30978.05 0.7573555 17721.23
##      93 31102.51 0.7575753 17796.24
##      94 31296.32 0.7581402 17934.13
##      95 31544.28 0.7573465 18033.62
##      96 31710.51 0.7579952 18127.32
##      97 31827.47 0.7552275 18141.07
##      98 31908.02 0.7541961 18164.55
##      99 32002.40 0.7526284 18200.99
##     100 32132.82 0.7500400 18254.19
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was k = 4.
```

The chosen model can then be used to predict just as before.

```
predkNN <- predict(kNNFit, newdata = test)
postResample(predkNN, test$selling_price)
```

```
##           RMSE      Rsquared      MAE
## 4.504931e+04 8.224044e-01 1.768072e+04
```

The same process can be used to fit and predict for an SLR or MLR model.

```
SLRFit <- train(selling_price ~ ex_showroom_price,
               data = train,
               method = "lm",
               trControl = trainControl(method = "cv", number = 10)
               )
SLRFit
```

```
## Linear Regression
##
## 438 samples
## 1 predictor
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 395, 394, 393, 394, 395, 394, ...
## Resampling results:
##
##      RMSE      Rsquared      MAE
## 21397.87 0.8006818 15128.29
##
## Tuning parameter 'intercept' was held constant at a value of TRUE
predSLR <- predict(SLRFit, newdata = test)
postResample(predSLR, test$selling_price)

##           RMSE      Rsquared      MAE
## 2.784060e+04 8.845011e-01 1.689184e+04
```

Multiple Predictors

Just like SLR can include multiple explanatory variables, we can include multiple explanatory variables with kNN (they must all be numeric unless you develop or use a ‘distance’ measure that is appropriate for categorical data).

With all numeric explanatory variables, we often use Euclidean distance as our distance metric. For instance, with two explanatory variables x_1 and x_2 :

$$d(x_1, x_2) = \sqrt{(x_{11} - x_{12})^2 + (x_{21} - x_{22})^2}$$

The same model notation from before can be used:

$$\text{response variable} \sim \text{explanatory variable1} + \text{explanatory variable2} + \dots$$

Along with the same kind of R code to fit the model:

```
k <- 1:100
kNNFit <- train(selling_price ~ ex_showroom_price + km_driven + year,
  data = train,
  method = "knn",
  tuneGrid = data.frame(k = k),
  trControl = trainControl(method = "cv", number = 10)
)
kNNFit
```

```
## k-Nearest Neighbors
##
## 438 samples
## 3 predictor
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 394, 394, 396, 394, 395, ...
## Resampling results across tuning parameters:
##
##  k    RMSE      Rsquared    MAE
##  1  23640.50  0.7573258  16460.14
##  2  21836.42  0.7857472  14837.83
##  3  21019.79  0.7938294  13982.71
##  4  20922.74  0.8033729  13642.88
##  5  21341.90  0.8057949  13772.31
##  6  21567.20  0.8101463  13833.95
##  7  21779.42  0.8111850  13944.40
##  8  21996.73  0.8120830  13937.33
##  9  22049.90  0.8107130  13867.25
## 10  21997.21  0.8132944  13782.32
## 11  22144.41  0.8109825  13807.93
## 12  22260.30  0.8090794  13795.21
## 13  22320.04  0.8104700  13740.63
## 14  22409.88  0.8108319  13750.92
## 15  22582.47  0.8074337  13847.22
## 16  22587.63  0.8073978  13802.44
## 17  22495.87  0.8096101  13765.12
## 18  22502.31  0.8097275  13758.29
## 19  22635.53  0.8080655  13750.78
```

##	20	22703.92	0.8063917	13735.66
##	21	22734.59	0.8063965	13705.90
##	22	22780.80	0.8055256	13706.06
##	23	22786.35	0.8063212	13692.50
##	24	22901.19	0.8046647	13756.52
##	25	23064.28	0.8030037	13840.15
##	26	22945.98	0.8057186	13810.70
##	27	23050.27	0.8033385	13845.49
##	28	23095.08	0.8033021	13836.57
##	29	23065.82	0.8038240	13818.44
##	30	23068.04	0.8043034	13809.20
##	31	22987.24	0.8058844	13817.12
##	32	22949.61	0.8065536	13841.69
##	33	22932.02	0.8076682	13844.54
##	34	22967.88	0.8075008	13866.61
##	35	22986.66	0.8070167	13913.72
##	36	23017.12	0.8068164	13940.93
##	37	23014.74	0.8072177	13935.62
##	38	23053.59	0.8069973	13955.12
##	39	23121.92	0.8058505	13976.88
##	40	23188.98	0.8054436	13985.18
##	41	23242.65	0.8044281	14024.84
##	42	23307.17	0.8038732	14041.14
##	43	23347.23	0.8034174	14046.71
##	44	23400.94	0.8030069	14088.13
##	45	23419.83	0.8037873	14093.62
##	46	23457.56	0.8045055	14118.09
##	47	23534.76	0.8044310	14155.88
##	48	23550.50	0.8055000	14164.25
##	49	23608.15	0.8060852	14190.60
##	50	23680.20	0.8061208	14218.41
##	51	23792.80	0.8055961	14248.96
##	52	23863.67	0.8061005	14237.46
##	53	23941.24	0.8057372	14268.26
##	54	24012.78	0.8064626	14283.52
##	55	24088.04	0.8066189	14314.38
##	56	24223.90	0.8058238	14380.16
##	57	24332.59	0.8057464	14429.05
##	58	24432.58	0.8056225	14483.27
##	59	24471.53	0.8072791	14499.97
##	60	24590.55	0.8072288	14528.39
##	61	24706.51	0.8068890	14549.87
##	62	24824.34	0.8063164	14603.47
##	63	24974.46	0.8056767	14662.43
##	64	25129.51	0.8052308	14721.67
##	65	25210.99	0.8059404	14761.42
##	66	25212.34	0.8079921	14782.19
##	67	25361.68	0.8069921	14852.17
##	68	25522.10	0.8053265	14928.02
##	69	25551.98	0.8074049	14938.67
##	70	25684.36	0.8067109	14999.57
##	71	25801.00	0.8067485	15048.63
##	72	25923.95	0.8063366	15110.98
##	73	26049.21	0.8058486	15179.51


```
##      74 26150.98 0.8067401 15223.84
##      75 26287.68 0.8064546 15284.45
##      76 26389.27 0.8070254 15340.69
##      77 26516.86 0.8066690 15397.88
##      78 26623.45 0.8070039 15458.81
##      79 26744.78 0.8069820 15507.65
##      80 26888.79 0.8060954 15576.07
##      81 27036.57 0.8051313 15656.25
##      82 27172.33 0.8047743 15716.55
##      83 27291.07 0.8049043 15774.40
##      84 27401.84 0.8049215 15823.50
##      85 27546.69 0.8047449 15907.97
##      86 27671.29 0.8047395 15966.98
##      87 27798.12 0.8042159 16043.73
##      88 27926.29 0.8039591 16102.04
##      89 28070.78 0.8032271 16188.30
##      90 28197.69 0.8026942 16240.57
##      91 28308.91 0.8023809 16305.39
##      92 28448.75 0.8020104 16380.77
##      93 28545.18 0.8022831 16424.13
##      94 28663.66 0.8023173 16485.83
##      95 28784.19 0.8023103 16538.29
##      96 28914.83 0.8019965 16600.57
##      97 29030.81 0.8015904 16665.41
##      98 29131.86 0.8007397 16711.59
##      99 29247.84 0.8006306 16767.39
##     100 29350.81 0.8008318 16827.80
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was k = 4.
```

The chosen model can then be used to predict just as before.

```
predkNN <- predict(kNNFit, newdata = test)
postResample(predkNN, test$selling_price)
```

```
##           RMSE      Rsquared      MAE
## 4.474815e+04 8.233754e-01 1.722657e+04
```

Just for reference: let's compare this to the MLR output.

```
MLRFit <- train(selling_price ~ ex_showroom_price + km_driven + year,
  data = train,
  method = "lm",
  trControl = trainControl(method = "cv", number = 10)
)
MLRFit
```

```
## Linear Regression
##
## 438 samples
## 3 predictor
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 394, 393, 396, 394, 394, 395, ...
## Resampling results:
```

```
##
##      RMSE      Rsquared  MAE
##    17256.27  0.877712  11293.37
##
## Tuning parameter 'intercept' was held constant at a value of TRUE
predMLR <- predict(MLRFit, newdata = test)
postResample(predMLR, test$selling_price)
```

```
##      RMSE      Rsquared      MAE
## 2.337461e+04 9.215234e-01 1.191222e+04
```

Note: Practical use of kNN says we should usually standardize (center to have mean 0 and scale to have standard deviation 1) our numeric explanatory variables. Why?

Day 5

Competition!

Time to put what we've learned into practice! [Kaggle](#) is a site that hosts competitions around predicting a response (either a numeric response or predicting the category that an observation might belong to).

Housing Prices

Let's go check out our competition: <https://www.kaggle.com/c/house-prices-advanced-regression-techniques/overview>

Use the starter files to come up with some models!