

Prediction!

Contents

Prediction	1
Relating Explanatory Variables in Prediction	4
Fitting a Linear Regression Model in R	8
Another Modeling Approach (k Nearest Neighbors)	18
Competition!	20

Ok, so I think I start off a bit loose to get people comfortable and bring out some playing cards.

- I'll point to someone and ask them to guess the suit of the next card I show.
- Reshuffle and repeat giving them five cards/guesses.
- Then I'll repeat with another three-four students.
- # correctly guessed will be noted somewhere
- They've done a similar thing before but now an emphasis on guessing the next # of correct.


What's the point? How can I best predict the number of card suits the next person will get right?

Prediction

Goal: Predict a new value of a variable

- Ex: Another student will be guessing. Define $Y = \#$ of card suits guessed correctly from the five. What should we guess/predict for the next value of Y ?

- Chat with them. - Lead them to talk about a sample. - Simulate values of Y using an app (in the repo, use `shiny::runGitHub("caryAcademy", username = "jbpost2", subdir = "CardSim", ref = "main")`). - Lead them to ideas of using something like the sample mean or median as the predicted value. - Why something like the sample mean or sample median? What are we really trying to do? Find a value that is 'close' to the most values, i.e. something in the center being the most logical thing to do.



Loss function

Let's assume we have a sample of n people that each guessed five cards. Call these values y_1, y_2, \dots, y_n .

Need: A way to quantify how well our prediction is doing... Suppose there is some best prediction, call it c . How do we measure the quality of c ?

- Using the idea that we want something ‘close’ to all points, we find a way to compare each point to our prediction. - Think about things like:

$$y_1 - c, (y_1 - c)^2, |y_1 - c|$$

$$\sum_{i=1}^n (y_i - c), \sum_{i=1}^n (y_i - c)^2, \sum_{i=1}^n |y_i - c|$$

$$\frac{1}{n} \sum_{i=1}^n (y_i - c), \frac{1}{n} \sum_{i=1}^n (y_i - c)^2, \frac{1}{n} \sum_{i=1}^n |y_i - c|$$

- Quick app to look at how the measures work in the app. - In the end an objective function (mean squared error here) must be created to minimize that uses a ‘Loss function’, and we’ll talk about why we’ll use the common squared error loss:

$$g(y_1, \dots, y_n) = \frac{1}{n} \sum_{i=1}^n L(y_i, c) = \frac{1}{n} \sum_{i=1}^n (y_i - c)^2$$

Can we choose an ‘optimal’ value for c to minimize this function? Calculus to the rescue!

Steps to minimize a function with respect to c :

1. Take the derivative with respect to c
2. Set the derivative equal to 0
3. Solve for c to obtain the potential maximum or minimum
4. Check to see if you have a maximum or minimum (or neither)

Answer comes out to be \bar{y} as the minimizer.

Big wrap: This means that the sample mean is the best prediction when using squared error loss (root mean square error).

Using a Population Distribution

Rather than using sample data, suppose we think about the theoretical distribution for $Y = \#$ of card suits guessed correctly from the five. What might we use here? What assumptions do we need to make this distribution reasonable?

- $Y \sim \text{Bin}(5, 0.25)$ assuming we have independent and identical trials - This gives

$$p_Y(y) = P(Y = y) = \binom{n}{y} p^y (1-p)^{n-y} = \binom{5}{y} 0.25^y 0.75^{5-y}$$

for $y = 0, 1, 2, \dots, n$ or $y = 0, 1, 2, 3, 4, 5$

Is there an optimal value c for the **expected value** of the loss function?

That is, can we minimize (as a function of c) $E[(Y - c)^2]$?

- I think I’ll start them out on this one as theory leads to more difficult ideas and I’m not sure if you did general expected values. - In the end though we end up with np or 1.25. - I’ll show/discuss that this works generally for any distribution $p_Y(y)$ that has a mean - Discuss the relationship with this and a sample (1/n weight for each point vs $p_Y(y)$ weight for each.) - **Big idea:** This implies that μ is the best predictor to use if you are considering minimizing the expected squared error loss.

HW for after day 1:

- Give them a data set and have them find the mean in R and note that the prediction they would use is that.
- Have them play minesweeper and record the data appropriately. Ask them to produce a best guess for their # of overall bombs.
- Give them a partial derivative question to practice on.

This would be the 2nd (short day) material.

- Recap the big idea of prediction:
 - Need to quantify how well we are doing (squared error loss and MSE)
 - Sample mean is optimal if we have a sample
 - Given a theoretical distribution, expected squared error loss is optimized at the mean of the distribution
- Introduce next material with minesweeper, because it is now browser based, nostalgia on my part, and it seems somewhat fun <https://minesweeper.online/game/938135731>
 - Each student will be assigned a certain number of mines for the board (15x40).
 - They'll click on the first square just below the smiley face. They'll continue to click down one block at a time until they hit a bomb.
 - They'll record in a shared spreadsheet their number of blocks down the first bomb appears.
 - Each person should play 10 games and put their data in.
- Now we'll discuss how we could predict the number of blocks until the first bomb as a function of the number of bombs.
- We'll read in the data to R and do some plotting (I'm not sure what the relationship will be exactly but I'd guess not super linear).

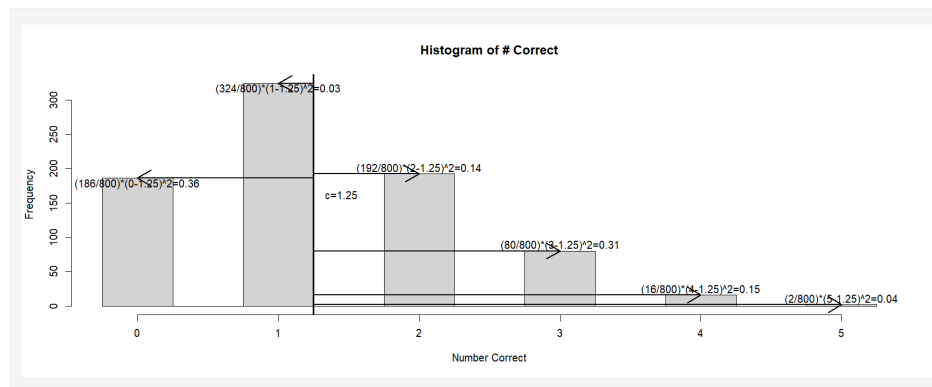
Relating Explanatory Variables in Prediction

Y is a random variable and we'll consider the x values fixed (we'll denote this as $Y|x$). We hope to learn about the relationship between Y and x .

When we considered just Y by itself and used squared error loss, we know that $E(Y) = \mu$ minimizes


$$E[(Y - c)^2]$$

as a function of c . Given data, we used $\hat{\mu} = \bar{y}$ as our prediction.



Harder (and more interesting) problem is to consider predicting a (response) variable Y as a function of an explanatory variable x .

Below: Blue line, $f(x)$, is the 'true' relationship between x and y



Now that we have an x , $E(Y|x)$ will minimize

$$E[(Y - c)^2|x]$$


We can call this true unknown value $E(Y|x) = f(x)$. That is, the average value of Y will now be considered as a function of x .

Given observed Y 's and x 's, we can estimate this function as $\hat{f}(x)$ (think \bar{y} from before). This $\hat{f}(x)$ will minimize


$$g(y_1, \dots, y_n | x_1, \dots, x_n) = \frac{1}{n} \sum_{i=1}^n L(y_i, \hat{f}(x_i)) = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{f}(x_i))^2$$

Approximating $f(x)$

Although the true relationship is most certainly nonlinear, we may be ok approximating the relationship linearly. For example, consider the same plot as above but between 0 and 5 only:



That's pretty linear. Consider plot between 5 and 15:



Line still does a reasonable job and is often used as a basic approximation.

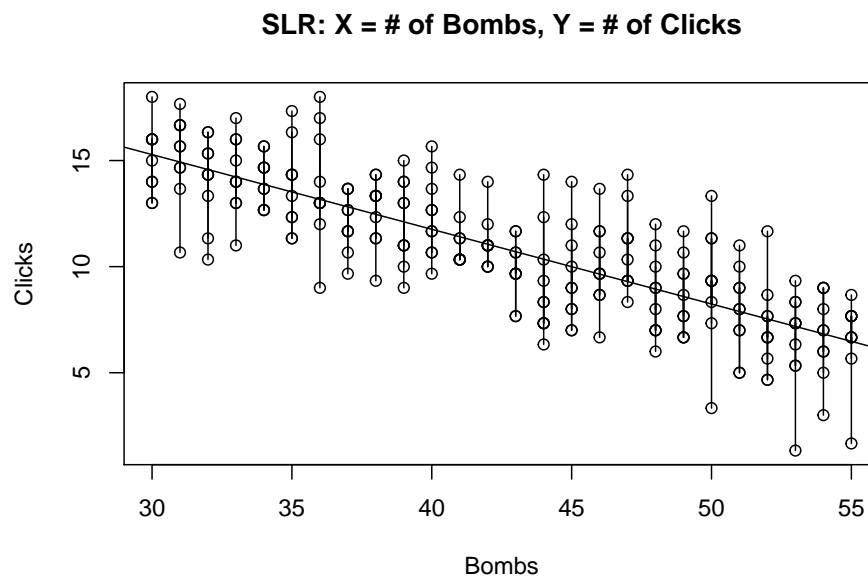
Linear Regression Model

The (fitted) linear regression model uses $\hat{f}(x) = \hat{\beta}_0 + \hat{\beta}_1 x$. This means we want to find the optimal values of $\hat{\beta}_0$ and $\hat{\beta}_1$ from:

$$g(y_1, \dots, y_n | x_1, \dots, x_n) = \sum_{i=1}^n (y_i - \hat{\beta}_0 - \hat{\beta}_1 x_i)^2$$

This equation is often called the 'sum of squared errors (or residuals)' or the 'residual sum of squares'. The model for the data, $E(Y|x) = f(x) = \beta_0 + \beta_1 x$ is called the Simple Linear Regression (SLR) model.

I'll have code at the ready to update and rerender this plot using their data from minesweeper.



Calculus allows us to find the ‘least squares’ estimators, $\hat{\beta}_0$ and $\hat{\beta}_1$ in a nice closed-form!

Do they know partial derivatives? I’m not sure. I think we’ll be running low on time here anyway, so maybe I’ll just talk about the idea of how to get them, set up the equations and then just give the answers.

$$\hat{\beta}_1 = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (x_i - \bar{x})^2}$$

$$\hat{\beta}_0 = \bar{y} - \bar{x}\hat{\beta}_1$$

Then we’ll jump into R and find the values for the minesweeper data and use it to predict by “hand” (plugging it in manually in R).

Ok, day 3 here. Get into R and do some model fitting and predicting. Start with a quick recap here

Fitting a Linear Regression Model in R

Recap: Our goal is to predict a value of Y while including an explanatory variable x . We are assuming we have a sample of (x_i, y_i) pairs, $i = 1, \dots, n$.

The Simple Linear Regression (SLR) model can be used:

$$\hat{f}(x_i) = \hat{y}_i = \hat{\beta}_0 + \hat{\beta}_1 x_i$$

where

- y_i is our response for the i^{th} observation
- x_i is the value of our explanatory variable for the i^{th} observation
- β_0 is the y intercept
- β_1 is the slope

The best model to use if we consider squared error loss has

$$\hat{\beta}_1 = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (x_i - \bar{x})^2}$$
$$\hat{\beta}_0 = \bar{y} - \bar{x}\hat{\beta}_1$$

called the ‘least squares estimates’.

Data Intro

This [dataset](#) contains information about used motorcycles and their cost.

From the information page: This data can be used for a lot of purposes such as price prediction to exemplify the use of linear regression in Machine Learning. The columns in the given dataset are as follows:

- name
- selling price
- year
- seller type
- owner
- km driven
- ex showroom price

The data are available to download from this URL:

<https://www4.stat.ncsu.edu/~online/datasets/bikeDetails.csv>

I think the exploration and modeling would be better to do live and ask them for input as we go through it. I'll put some stuff here and then we can talk about it.

Read in Data and Explore!

```
library(tidyverse)
bikeData <- read_csv("https://www4.stat.ncsu.edu/~online/datasets/bikeDetails.csv")
bikeData <- bikeData %>% tidyr::drop_na()
select(bikeData, selling_price, year, km_driven, ex_showroom_price, name, everything())
```



```
## # A tibble: 626 x 7
##   selling_price year km_driven ex_showroom_price name seller_type owner
##   <dbl> <dbl> <dbl> <dbl> <chr> <chr> <chr>
## 1 150000 2018 12000 148114 Royal Enfi~ Individual 1st ~
## 2 65000 2015 23000 89643 Yamaha Faz~ Individual 1st ~
## 3 18000 2010 60000 53857 Honda CB T~ Individual 1st ~
## 4 78500 2018 17000 87719 Honda CB H~ Individual 1st ~
## 5 50000 2016 42000 60122 Bajaj Disc~ Individual 1st ~
## 6 35000 2015 32000 78712 Yamaha FZ16 Individual 1st ~
## 7 28000 2016 10000 47255 Honda Navi Individual 2nd ~
## 8 80000 2018 21178 95955 Bajaj Aven~ Individual 1st ~
## 9 365000 2019 1127 351680 Yamaha YZF~ Individual 1st ~
## 10 25000 2012 55000 58314 Suzuki Acc~ Individual 1st ~
## # ... with 616 more rows
```

Our 'response' variable here is the `selling_price` and we could use the variable `year`, `km_driven`, or `ex_showroom_price` as the explanatory variable. Let's make some plots and summaries to explore.

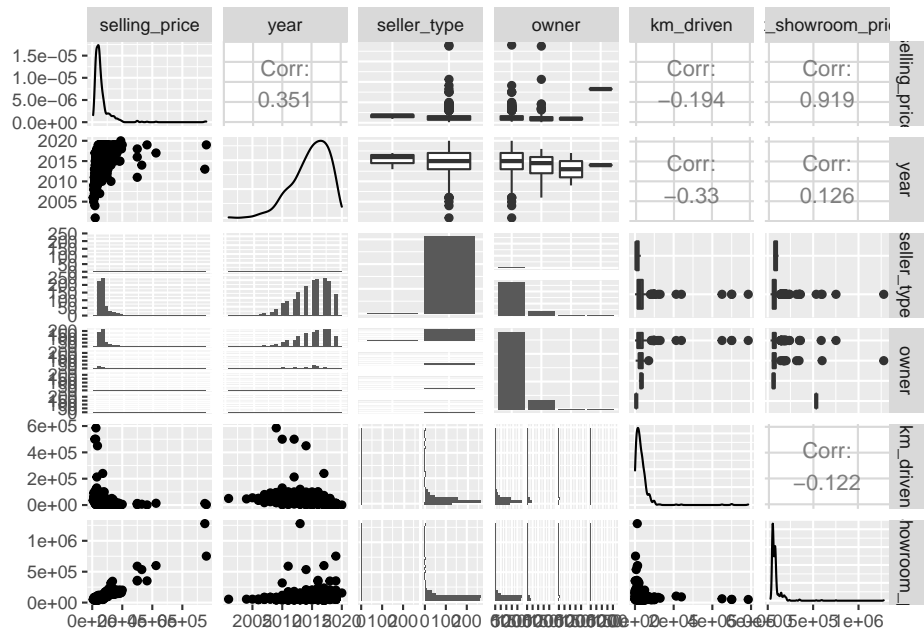
```
summary(bikeData)
```

```
##      name      selling_price      year      seller_type
## Length:626    Min.   : 6000    Min.   :2001    Length:626
## Class :character 1st Qu.: 30000 1st Qu.:2013    Class :character
## Mode :character  Median : 45000  Median :2015    Mode :character
##                Mean   : 59445   Mean   :2015
##                3rd Qu.: 65000   3rd Qu.:2017
##                Max.   :760000   Max.   :2020
##      owner      km_driven      ex_showroom_price
## Length:626    Min.   : 380    Min.   : 30490
## Class :character 1st Qu.: 13031 1st Qu.: 54852
## Mode :character  Median : 25000  Median : 72753
##                Mean   : 32672   Mean   : 87959
##                3rd Qu.: 40000   3rd Qu.: 87032
##                Max.   :585659   Max.   :1278000
```


```
summarize(group_by(bikeData, owner),
  mean = mean(selling_price),
  median = median(selling_price),
  sd = sd(selling_price),
  IQR = IQR(selling_price))
```

```
## # A tibble: 4 x 5
##   owner      mean median      sd  IQR
##   <chr>    <dbl> <dbl> <dbl> <dbl>
## 1 1st owner 58432. 45000 51125. 35000
## 2 2nd owner 64795. 35000 104861. 30750
## 3 3rd owner 39333. 40000 17010. 17000
## 4 4th owner 330000 330000    NA      0
```

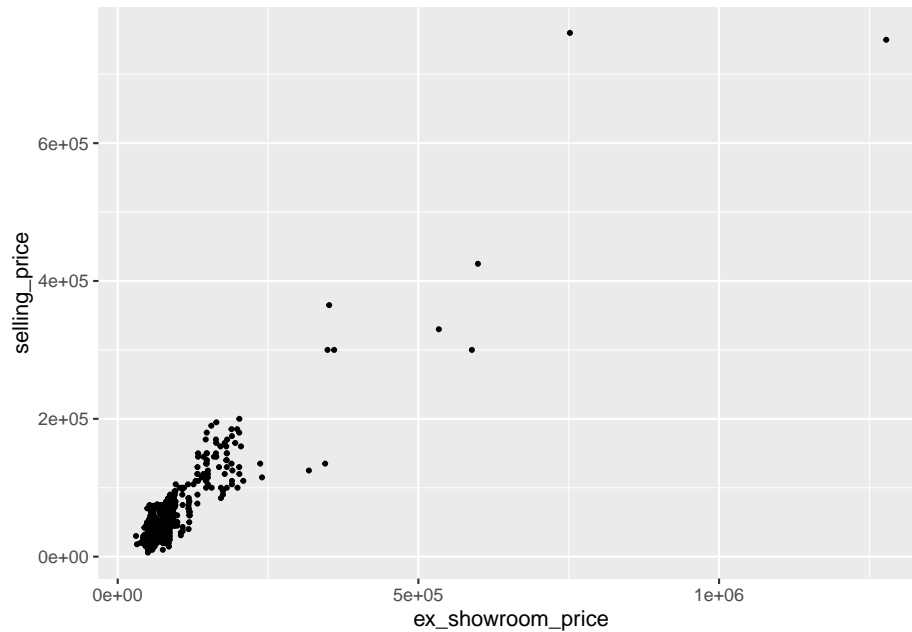
```
library(GGally)
ggpairs(select(bikeData, -name))
```




```
g <- ggplot(data = bikeData, aes(y = selling_price))
g + geom_jitter(aes(x = year, color = km_driven))
```



```
g + geom_point(aes(x = ex_showroom_price), size = 0.75)
```



```
g <- ggplot(data = filter(bikeData, ex_showroom_price < 100000), aes(y = selling_price))
g +
  geom_point(aes(x = ex_showroom_price, color = year, size = km_driven)) +
  scale_color_gradientn(colours = rainbow(5))
```



‘Fitting’ the Model

Basic *linear model* fits done with `lm()`. First argument is a **formula**:

$$\text{response variable} \sim \text{modeling variable}(s)$$

We specify the modeling variable(s) with a `+` sign separating variables. With SLR, we only have one variable on the right hand side.

```
fit <- lm(selling_price ~ ex_showroom_price, data = bikeData)
fit

##
## Call:
## lm(formula = selling_price ~ ex_showroom_price, data = bikeData)
##
## Coefficients:
##      (Intercept)  ex_showroom_price
##      -3010.6984         0.7101
```

We can easily pull off things like the coefficients.

```
coefficients(fit) #helper function

##      (Intercept) ex_showroom_price
##      -3010.6984021         0.7100588
```


Manually predict for an ex_showroom_price of 50000:

```
intercept <- coefficients(fit)[1]
slope <- coefficients(fit)[2]
intercept + slope * 50000
```

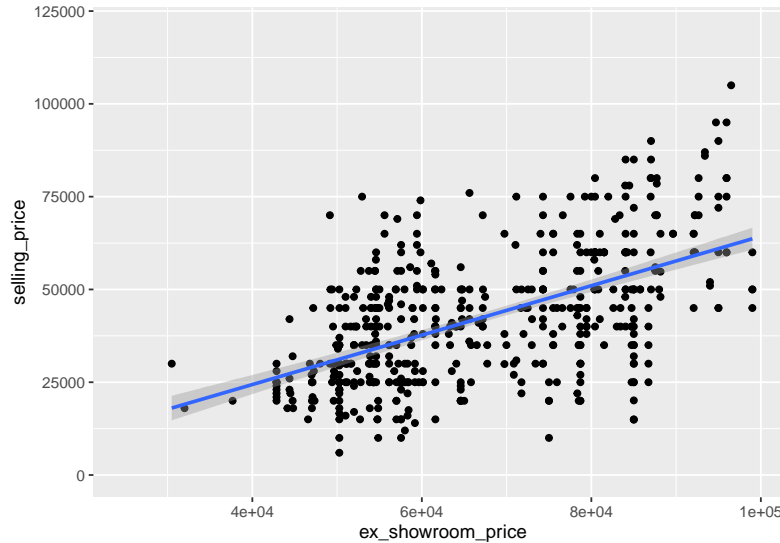
```
## (Intercept)
##      32492.24
```

We can also look at the fit of the line on the graph.

```
ggplot(bikeData, aes(x = ex_showroom_price, y = selling_price)) +
  geom_point() +
  geom_smooth(method = "lm")
```



```
ggplot(bikeData, aes(x = ex_showroom_price, y = selling_price)) +
  geom_point() +
  geom_smooth(method = "lm") +
  scale_x_continuous(limits = c(25000, 100000)) +
  scale_y_continuous(limits = c(0, 120000))
```



Predicting!

Can predict the `selling_price` for a given `ex_showroom_price` easily using the `predict()` function.

```
predict(fit, newdata = data.frame(ex_showroom_price = c(50000, 75000, 100000)))
```

```
##          1          2          3
## 32492.24 50243.71 67995.19
```

Error Assumptions

Although, not needed for prediction, we often assume that we observe our response variable Y as a function of the line plus random errors:

$$Y_i = \beta_0 + \beta_1 x_i + E_i$$

where the errors come from a Normal distribution with mean 0 and variance σ^2 ($E_i \stackrel{iid}{\sim} N(0, \sigma^2)$)

If we do this and use probability theory (maximum likelihood), we will get the same estimates for the slope and intercept as above!

What we get from the normality assumption (if reasonable) is the knowledge of the distribution of our estimators ($\hat{\beta}_0$ and $\hat{\beta}_1$).

What does knowing the distribution allow us to do? We can create confidence intervals or conduct hypothesis tests.

- Get standard error (SE) for prediction

```
predict(fit, newdata = data.frame(ex_showroom_price = c(50000, 75000, 100000)), se.fit = TRUE)
```

```
## $fit
##          1          2          3
## 32492.24 50243.71 67995.19
##
## $se.fit
##          1          2          3
## 1054.7005  960.2046  958.4166
##
## $df
```

```
## [1] 624
##
## $residual.scale
## [1] 23694.8
```

- Get confidence interval for mean response

```
predict(fit, newdata = data.frame(ex_showroom_price = c(50000, 75000, 100000)),
       se.fit = TRUE, interval = "confidence")
```

```
## $fit
##      fit      lwr      upr
## 1 32492.24 30421.05 34563.44
## 2 50243.71 48358.09 52129.34
## 3 67995.19 66113.07 69877.30
##
## $se.fit
##      1      2      3
## 1054.7005 960.2046 958.4166
##
## $df
## [1] 624
##
## $residual.scale
## [1] 23694.8
```

- Get prediction interval for new response


```
predict(fit, newdata = data.frame(ex_showroom_price = c(50000, 75000, 100000)),
       se.fit = TRUE, interval = "prediction")
```

```
## $fit
##      fit      lwr      upr
## 1 32492.24 -14085.045 79069.53
## 2 50243.71  3674.309 96813.12
## 3 67995.19 21425.922 114564.45
##
## $se.fit
##      1      2      3
## 1054.7005 960.2046 958.4166
##
## $df
## [1] 624
##
## $residual.scale
## [1] 23694.8
```

- Can see the confidence and prediction bands on the plot:

```
library(ciTools)
bikeData <- add_pi(bikeData, fit, names = c("lower", "upper"))

ggplot(bikeData, aes(x = ex_showroom_price, y = selling_price)) +
  geom_point() +
  geom_smooth(method = "lm", fill = "Blue") +
  geom_ribbon(aes(ymin = lower, ymax = upper), alpha = 0.3, fill = "Red") +
  ggtitle("Scatter Plot with 95% PI & 95% CI")
```



For HW have them jump into R and run the code to read in the minesweeper data. Then they could use `lm()` to fit a model and predict.

Multiple Linear Regression

We can add in more than one explanatory variable using the formula for `lm()`. The ideas all follow through!

```
fit <- lm(selling_price ~ ex_showroom_price + year + km_driven, data = bikeData)
fit
```

```
##
## Call:
## lm(formula = selling_price ~ ex_showroom_price + year + km_driven,
##     data = bikeData)
##
## Coefficients:
##      (Intercept)  ex_showroom_price          year      km_driven
##      -9.429e+06      6.863e-01      4.679e+03     -1.053e-02
```

To predict we now need to specify values for all the explanatory variables.

```
data.frame(ex_showroom_price = c(50000, 75000),
           year = c(2010, 2011),
           km_driven = c(15000, 10000))
```

```
##   ex_showroom_price year km_driven
## 1          50000 2010    15000
## 2          75000 2011    10000
```

```
predict(fit, newdata = data.frame(ex_showroom_price = c(50000, 75000),
                                   year = c(2010, 2011),
                                   km_driven = c(15000, 10000)),
       se.fit = TRUE, interval = "confidence")
```

```
## $fit
```

```
##           fit           lwr           upr
## 1 11118.83  7914.815 14322.85
## 2 33007.56 30202.482 35812.63
##
## $se.fit
##           1           2
## 1631.552 1428.402
##
## $df
## [1] 622
##
## $residual.scale
## [1] 19011.31
```

Difficult to visualize the model fit though!

Evaluating Model Accuracy

Which model is better? Ideally we want a model that can predict **new** data better, not the data we've already seen. We need a **test** set to predict on. We also need to quantify what we mean by better!

Training and Test Sets

We can split the data into a **training set** and **test set**.



- On the training set we can fit (or train) our models. The data from the test set isn't used at all in this process.
- We can then predict for the test set observations (for the combinations of explanatory variables seen in the test set). Can then compare the predicted values to the actual observed responses from the test set.

Split data randomly:

```
set.seed(1)
numObs <- nrow(bikeData)
index <- sample(1:numObs, size = 0.7*numObs, replace = FALSE)
train <- bikeData[index, ]
test <- bikeData[-index, ]
```

Fit the models on the training data only.

```
fitSLR <- lm(selling_price ~ ex_showroom_price , data = train)
fitMLR <- lm(selling_price ~ ex_showroom_price + year + km_driven, data = train)
```

Predict on the test set.

```
predSLR <- predict(fitSLR, newdata = test)
predMLR <- predict(fitMLR, newdata = test)
tibble(predSLR, predMLR, test$selling_price)
```



```
## # A tibble: 188 x 3
##   predSLR predMLR `test$selling_price`
##   <dbl>   <dbl>           <dbl>
## 1 100749. 113099.           150000
## 2  59739.  60915.             65000
## 3  58390.  72928.             78500
## 4  39035.  45467.             50000
## 5  52073.  53538.             35000
## 6  64166.  78343.             80000
## 7  79576.  57387.             40000
## 8 100749. 112955.           150000
## 9  89924. 102497.           120000
## 10 36247.  25302.             25000
## # ... with 178 more rows
```

Root Mean Square Error

Which is better?? Can use squared error loss to evaluate! (Square root of the mean squared error loss is often reported instead and is called RMSE or Root Mean Square Error.)

```
sqrt(mean((predSLR - test$selling_price)^2))
```

```
## [1] 27840.6
```

```
sqrt(mean((predMLR - test$selling_price)^2))
```

```
## [1] 23374.61
```

MLR fit does much better at predicting!

Another Modeling Approach (k Nearest Neighbors)

Day 4:

kNN relate to idea with minesweeper data about how they predicted for their individual piece.

Will need to use $k = 10, 30, 50, \dots$

Intro and Recap

Recap: Our previous goal was to predict a value of Y while including an explanatory variable x . With that x , we said $E(Y|x)$ will minimize

$$E[(Y - c)^2|x]$$

We called this true unknown value $E(Y|x) = f(x)$.

Given observed Y 's and x 's, we can estimate this function as $\hat{f}(x)$ (with SLR we estimated it with $\hat{\beta}_0 + \hat{\beta}_1 x$). This $\hat{f}(x)$ will minimize

$$g(y_1, \dots, y_n | x_1, \dots, x_n) = \frac{1}{n} \sum_{i=1}^n L(y_i, \hat{f}(x_i)) = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{f}(x_i))^2$$

What other things could we consider for $f(x)$???

Recall the minesweeper data we collected. In your homework, you were asked to determine a prediction of the number of blocks you could click before hitting a bomb for **your** given number of bombs. You were each creating your estimate of $f(x)$ for your x value!

Let's visualize that idea and compare it to the SLR fit!

I'd need the data here but this is what it would look like. I won't show the code here.

```
# means <- cars %>%
#   group_by(speed) %>%
#   summarize(mean = mean(dist))
# means
#
# #add first and last values to get break points for plotting
# library(zoo)
# speedBars <- c(min(cars$speed)-0.5, rollmean(unique(cars$speed), 2), max(cars$speed)+0.5)
# plot(x = cars$speed, y = cars$dist, main = "Using Local Mean as an Estimate", xlab = "Speed", ylab = "Distance")
# segments(x0 = speedBars[-length(speedBars)], x1 = speedBars[-1], y0 = means$mean, y1 = means$mean, lwd = 2)
#
# #add linear fit for comparison
# fit <- lm(dist ~ speed, data = cars)
# plot(x = cars$speed, y = cars$dist, main = "Using Local Mean as an Estimate", xlab = "Speed", ylab = "Distance")
# segments(x0 = speedBars[-length(speedBars)], x1 = speedBars[-1], y0 = means$mean, y1 = means$mean, lwd = 2)
# abline(fit, lwd = 2)

#get means
means <- minesweeper %>%
  group_by(bombs) %>%
  summarize(mean = mean(clicks))


## `summarise()` ungrouping output (override with `.groups` argument)
means

## # A tibble: 26 x 2
##   bombs mean
```

```
##      <int> <dbl>
## 1      30  15.1
## 2      31  15.3
## 3      32  14.1
## 4      33  14.3
## 5      34  14.1
## 6      35  13.7
## 7      36  13.8
## 8      37  12.2
## 9      38  12.6
## 10     39  12.1
## # ... with 16 more rows
```

```
#Local min vis
```


```
plot(x = minesweeper$bombs, y= minesweeper$clicks, main = "Using Local Mean", xlab = "Bombs", ylab = "Clicks",
      segments(x0 = unique(bombs)-0.5, x1 = unique(bombs)+0.5, y0 = means$mean, y1 = means$mean, lwd = 2))
```



```
#Compare to SLR
```

```
fit <- lm(clicks ~ bombs, data = minesweeper)
plot(x = minesweeper$bombs, y= minesweeper$clicks, main = "Using Local Mean vs SLR", xlab = "Bombs", ylab = "Clicks",
      abline(fit, lwd = 2)
      segments(x0 = unique(bombs)-0.5, x1 = unique(bombs)+0.5, y0 = means$mean, y1 = means$mean, lwd = 2))
```

Using Local Mean vs SLR



This is the idea of k Nearest Neighbors (kNN) for predicting a numeric response!

kNN

To predict a value of our (numeric) response kNN uses the **average of the k 'closest' responses**. For numeric data, we usually use Euclidean distance ($d(x_1, x_2) = \sqrt{(x_1 - x_2)^2}$) to determine the closest values.

- Large k implies more rigid (possibly *underfit* but lower variance prediction).
- Smaller k implies less rigid (possible *overfit* with high variance in prediction)

For the minesweeper data, we had many values at the same x (# of bombs). That's why we considered using only 10, 30, 50, ... Otherwise, we have ties and then things get tricky!

Practical use of kNN says we should standardize (center to have mean 0 and scale to have standard deviation 1) our numeric explanatory variables. Why?

Choosing the Value of k

Can do training and test or CV to determine k

Competition!

Last day would be discussion the Kaggle competition here: <https://www.kaggle.com/c/house-prices-advanced-regression-techniques/overview>

Perhaps split them into teams of two or three. Have them fit some models and submit to see how they do.

I'd have some basic code they could use to read the data in, split it into training and test so they can select their model, create predictions, and output them to a csv for submission (evaluation is log prediction minus log of the actual value).