

Data and Modeling

What makes something a statistical model?

What is the difference between prediction and inference?

Data


- When modeling, what should our data look like?

Relating Explanatory Variables to a Response Variable

Consider the response Y as a random variable. We'll consider the x values fixed (for any explanatory variable). Our interest is in learning about the relationship between Y and x .

Y is random, so we don't have a **deterministic** relationship...

Below: Blue line, $f(x)$, is the 'true' relationship between x and y




What should we try to relate/model?


Approximating $f(x)$

Although the true relationship is most certainly nonlinear, we may be ok approximating the relationship linearly. For example, consider the same plot as above but between 0 and 5 only:

Blue line, $f(x)$, is the 'true' relationship between x and y



That's pretty linear. Consider plot between 5 and 15:



Line still does a reasonable job and is often used as a basic approximation.

Exploratory Data Analysis (EDA)

What are our first steps with data?

Common steps to EDA

- 1.
- 2.
- 3.
- 4.
- 5.
- 6.

Data Intro

This [dataset](https://www4.stat.ncsu.edu/~online/datasets/bikeDetails.csv) contains information about used motorcycles and their cost.

From the information page: This data can be used for a lot of purposes such as price prediction to exemplify the use of linear regression in Machine Learning. The columns in the given dataset are as follows:

- name
- selling price
- year
- seller type
- owner
- km driven
- ex showroom price

The data are available to download from this URL:

<https://www4.stat.ncsu.edu/~online/datasets/bikeDetails.csv>

Read in Data and Explore!

```
library(tidyverse)
bikeData <- read_csv("https://www4.stat.ncsu.edu/~online/datasets/bikeDetails.csv")
select(bikeData, selling_price, year, km_driven, ex_showroom_price, name, everything())
```

```
## # A tibble: 1,061 x 7
##   selling_price year km_driven ex_showroom_price name          seller_type owner
##   <dbl> <dbl>    <dbl>          <dbl> <chr>          <chr>    <chr>
## 1      175000 2019      350              NA Royal Enfi~ Individual 1st ~
## 2       45000 2017     5650              NA Honda Dio  Individual 1st ~
## 3      150000 2018    12000      148114 Royal Enfi~ Individual 1st ~
## 4       65000 2015    23000      89643 Yamaha Faz~ Individual 1st ~
## 5       20000 2011    21000              NA Yamaha SZ ~ Individual 2nd ~
## 6       18000 2010   60000      53857 Honda CB T~ Individual 1st ~
## 7       78500 2018    17000      87719 Honda CB H~ Individual 1st ~
## 8      180000 2008    39000              NA Royal Enfi~ Individual 2nd ~
## 9       30000 2010    32000              NA Hero Honda~ Individual 1st ~
## 10      50000 2016    42000      60122 Bajaj Disc~ Individual 1st ~
## # ... with 1,051 more rows
```

Our 'response' variable here is the `selling_price` and we could use the variable `year`, `km_driven`, or `ex_showroom_price` as the explanatory variable. Let's make some plots and summaries to explore.

Linear Regression

Recap: Our goal is to predict a value of Y while including an explanatory variable x . We are assuming we have a sample of (x_i, y_i) pairs, $i = 1, \dots, n$.

The Simple Linear Regression (SLR) model can be used:

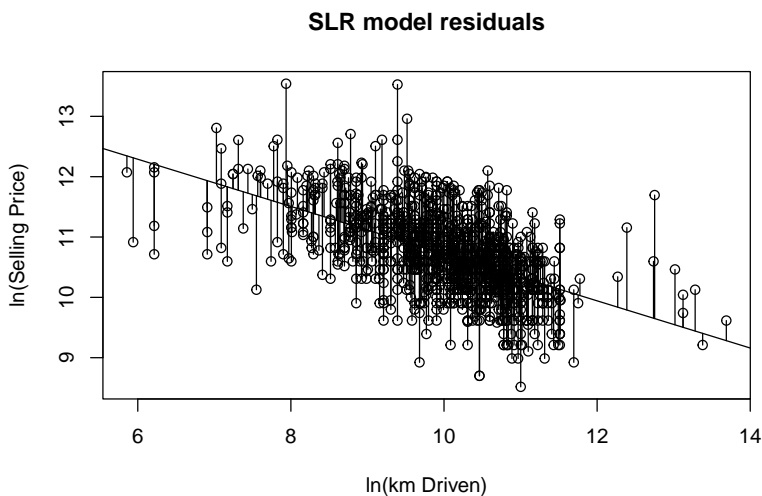
$$Y_i = \beta_0 + \beta_1 x_i + E_i$$

where

- y_i is our response for the i^{th} observation
- x_i is the value of our explanatory variable for the i^{th} observation
- β_0 is the y intercept
- β_1 is the slope
- $E_i \stackrel{iid}{\sim} N(0, \sigma^2)$

What is important to know from all that??

We **fit** this model to data. That is, find the **best** estimators of β_0 and β_1 (and σ^2) given the data. How to fit the line?



Fitting the line

The (fitted) linear regression model uses $\hat{f}(x) = \hat{\beta}_0 + \hat{\beta}_1 x$. Calculus allows us to find the ‘least squares’ estimators, $\hat{\beta}_0$ and $\hat{\beta}_1$ in a nice closed-form!

Making Inference

What hypothesis are we interested in and why?

How can we form a confidence interval for the quantity of interest?

Checking assumptions

How can we check our assumptions on the errors?

Fitting a Linear Regression Model in R

We can fit the model with the `lm()` function. Provide a formula

response ~ explanatory_variable_equation (intercept fit by default)

```
library(tidyverse)
bikeData <- read_csv("https://www4.stat.ncsu.edu/~online/datasets/bikeDetails.csv")
bikeData <- bikeData %>% mutate(log_selling_price = log(selling_price),
                                log_km_driven = log(km_driven))

fit <- lm(log_selling_price ~ log_km_driven, data = bikeData)
```

Determine the fitted model by looking at the `coefficients` element.

```
fit$coefficients
```

```
##      (Intercept) log_km_driven
##      14.6355683    -0.3910865
```

Look at the hypothesis test of interest with `summary()`

```
summary(fit)
```

```
##
## Call:
## lm(formula = log_selling_price ~ log_km_driven, data = bikeData)
##
```

```
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.9271 -0.3822 -0.0337  0.3794  2.5656
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  14.63557    0.18455   79.31  <2e-16 ***
## log_km_driven -0.39109    0.01837  -21.29  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.5953 on 1059 degrees of freedom
## Multiple R-squared:  0.2997, Adjusted R-squared:  0.299
## F-statistic: 453.2 on 1 and 1059 DF,  p-value: < 2.2e-16
```

What here is important and why?

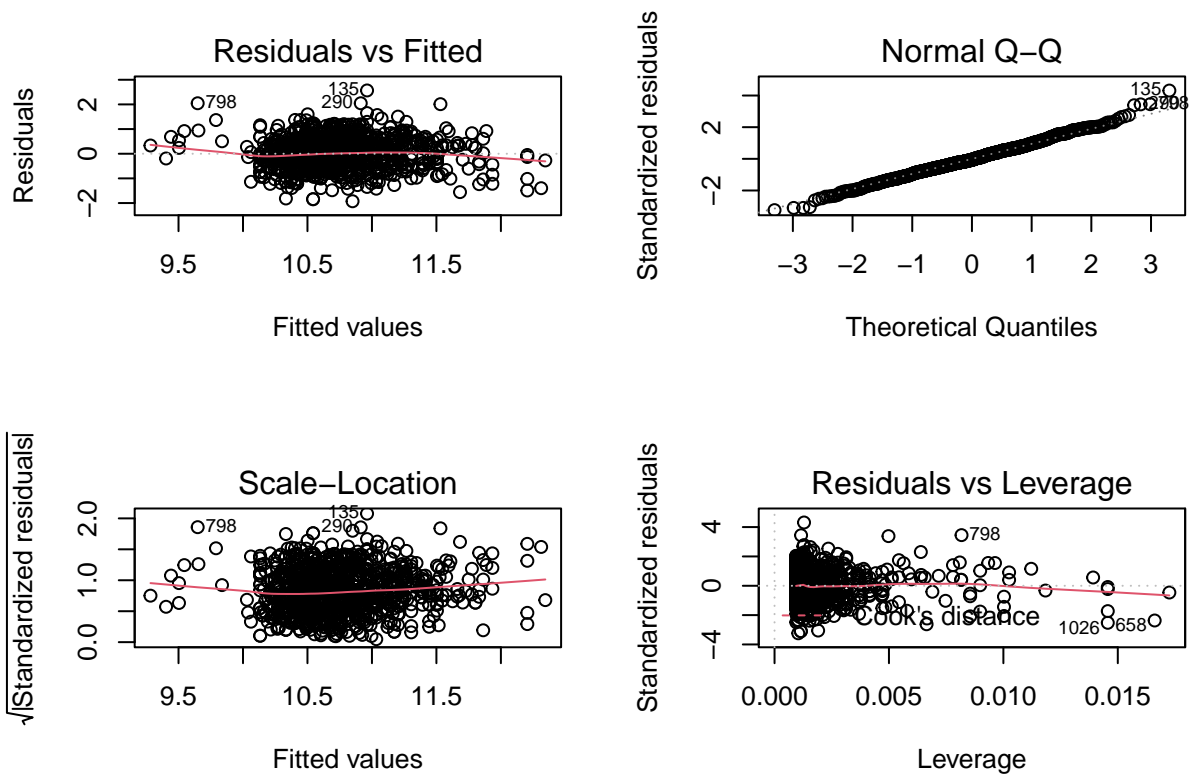
Find a confidence interval with `confint()`

```
confint(fit)
```

```
##              2.5 %      97.5 %
## (Intercept)  14.2734501 14.9976864
## log_km_driven -0.4271342 -0.3550389
```


Check conditions! `plot()` on the model fit will work.

```
par(mfrow = c(2,2))  
plot(fit)
```



Logistic Regression Model

Used when you have a **binary** response variable

- Using SLR is not appropriate!

Example:

- Consider data about [water potability](#)

```
library(tidyverse)
water <- read_csv("water_potability.csv")
water

## # A tibble: 3,276 x 10
##       ph Hardness Solids Chloramines Sulfate Conductivity Organic_carbon
##   <dbl>   <dbl>  <dbl>      <dbl>   <dbl>      <dbl>      <dbl>
## 1 NA      205.  20791.      7.30    369.      564.      10.4
## 2 3.72    129.  18630.      6.64     NA      593.      15.2
## 3 8.10    224.  19910.      9.28     NA      419.      16.9
## 4 8.32    214.  22018.      8.06    357.      363.      18.4
## 5 9.09    181.  17979.      6.55    310.      398.      11.6
## 6 5.58    188.  28749.      7.54    327.      280.       8.40
## 7 10.2    248.  28750.      7.51    394.      284.      13.8
## 8 8.64    203.  13672.      4.56    303.      475.      12.4
## 9 NA      119.  14286.      7.80    269.      389.      12.7
## 10 11.2    227.  25485.      9.08    404.      564.      17.9
## # ... with 3,266 more rows, and 3 more variables: Trihalomethanes <dbl>,
## #   Turbidity <dbl>, Potability <dbl>
```

- Summarize water potability

```
table(water$Potability)
```

```
##
##    0    1
## 1998 1278
```

```
water %>%
  group_by(Potability) %>%
  select(Hardness, Chloramines) %>%
  summary()
```


```
## Adding missing grouping variables: 'Potability'
```

```
##   Potability      Hardness      Chloramines
##   Min.   :0.0000   Min.    : 47.43   Min.    : 0.352
##   1st Qu.:0.0000   1st Qu.:176.85   1st Qu.: 6.127
##   Median :0.0000   Median :196.97   Median : 7.130
##   Mean   :0.3901   Mean    :196.37   Mean    : 7.122
##   3rd Qu.:1.0000   3rd Qu.:216.67   3rd Qu.: 8.115
##   Max.   :1.0000   Max.    :323.12   Max.    :13.127
```

Why is linear regression not appropriate?

```
fit <- lm(Potability ~ Hardness, data = water)
ggplot(water, aes(x = Hardness, y = Potability)) +
  geom_point() +
  geom_smooth(method = "lm")
```


'geom_smooth()' using formula 'y ~ x'



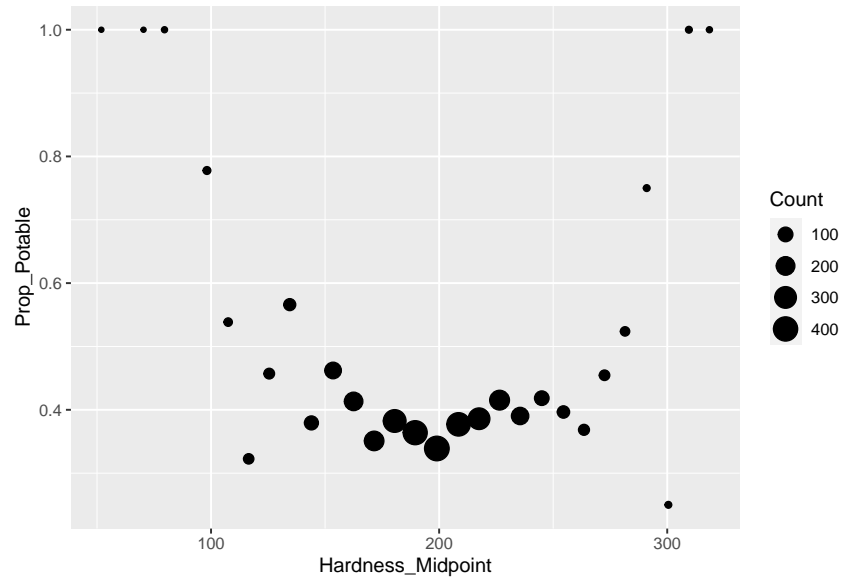
Better view...

```
fit <- lm(Potability ~ Hardness, data = water)
ggplot(water, aes(x = Hardness, y = Potability)) +
  geom_jitter() +
  geom_smooth(method = "lm")
```

'geom_smooth()' using formula 'y ~ x'



A better view of the data is to visualize the proportions of successes as a function of hardness.




- In SLR, we modeled the average of the response as a linear function. What does the average of the responses represent here? Why does using a linear function not make sense?

- Basic Logistic Regression models success probability using the *logistic function*

$$P(\text{success}|x) = \frac{e^{\beta_0 + \beta_1 x}}{1 + e^{\beta_0 + \beta_1 x}}$$

- This function never goes below 0 and never above 1 - works great for many applications!
- The logistic regression model doesn't have a closed form solution (maximum likelihood often used to fit parameters)



- Back-solving shows the *logit* or *log-odds* of success is linear in the parameters

$$\log \left(\frac{P(\text{success}|x)}{1 - P(\text{success}|x)} \right) = \beta_0 + \beta_1 x$$

- Coefficient interpretation changes greatly from linear regression model!
- β_1 represents a change in the log-odds of success

Hypotheses of Interest

What do you think would indicate that x is related to the probability of success here?

Fitting a Logistic Regression Model in R

Fit in R using `glm()` with `family = binomial` and a formula just like `lm()`.

```
fit <- glm(Potability ~ Hardness, data = water, family = "binomial")
```

Get coefficients by looking at `coefficients` element:

```
fit$coefficients
```

```
##      (Intercept)      Hardness  
## -0.2774792831 -0.0008629619
```

Get hypothesis test via `summary()`:

```
summary(fit)
```

```
##  
## Call:  
## glm(formula = Potability ~ Hardness, family = "binomial", data = water)  
##  
## Deviance Residuals:  
##      Min       1Q   Median       3Q      Max  
## -1.0279  -0.9963  -0.9853   1.3678   1.4209  
##  
## Coefficients:  
##              Estimate Std. Error z value Pr(>|z|)  
## (Intercept) -0.277479   0.216758  -1.280   0.200  
## Hardness    -0.000863   0.001090  -0.792   0.428  
##  
## (Dispersion parameter for binomial family taken to be 1)  
##  
##      Null deviance: 4382.0  on 3275  degrees of freedom  
## Residual deviance: 4381.3  on 3274  degrees of freedom  
## AIC: 4385.3  
##  
## Number of Fisher Scoring iterations: 4
```

Get confidence interval for β_1 with:

```
confint(fit)
```

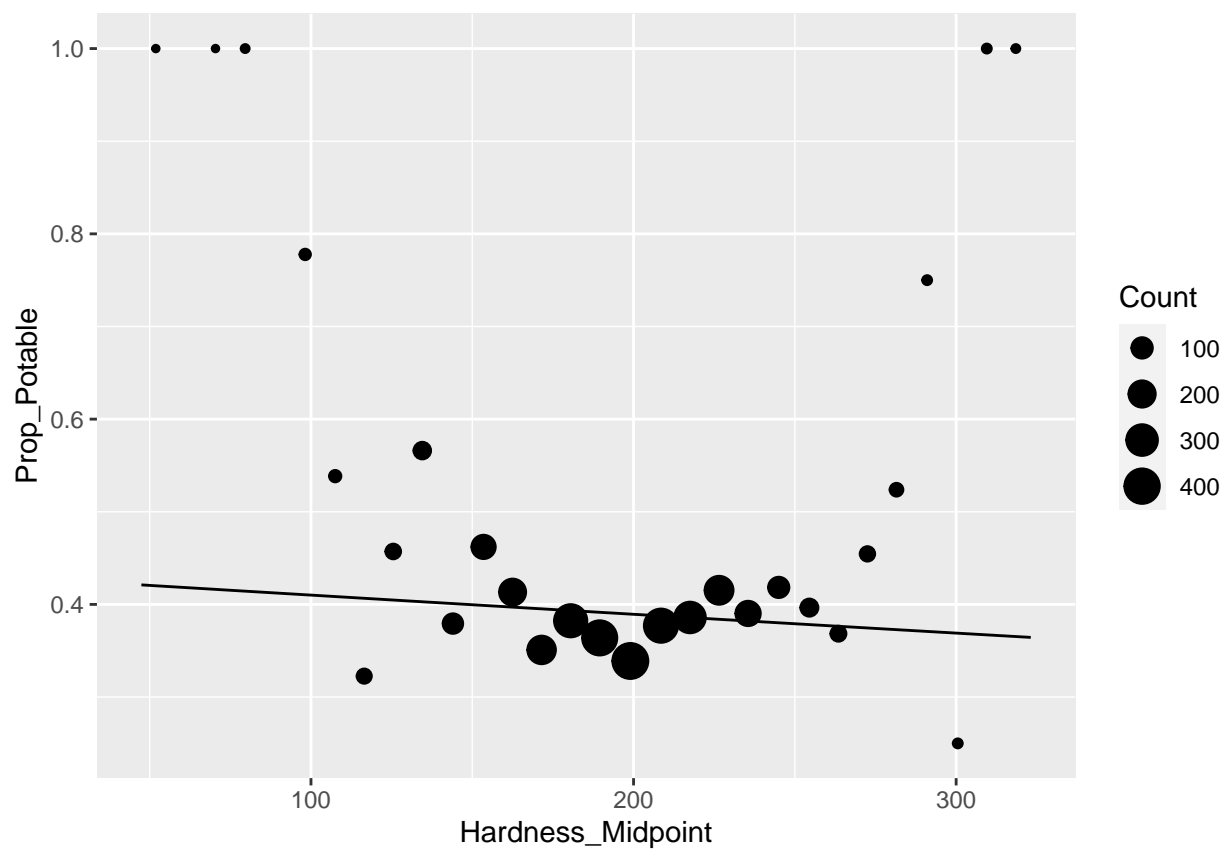
```
## Waiting for profiling to be done...  
  
##              2.5 %      97.5 %  
## (Intercept) -0.702803063 0.147169863  
## Hardness    -0.003000628 0.001272738
```


If we want a probability estimate back, use `predict()` with `type = 'link'`:

```
predict(fit, newdata = data.frame(Hardness = c(200, 300)), type = "link", se.fit = TRUE)
```

```
## $fit
##      1      2
## -0.4500717 -0.5363679
##
## $se.fit
##      1      2
## 0.03606504 0.11869267
##
## $residual.scale
## [1] 1
```

Visualize the fit:





Is a logistic curve!